

Keras -- MLPs on MNIST

In [80]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

In [81]:

```
from keras.initializers import glorot_normal
from keras.initializers import he_normal
```

In [82]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [83]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [84]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of test examples : 10000 and each image is of shape (28, 28)

In [85]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [86]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of test examples : 10000 and each image is of shape (784)

number of test examples : 10000 and each image is of shape (784)

In [87]:

```
# An example data point
print(X_train[0])
```

| | | | | | | | | | | | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|---|
| [| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 3 | 18 | 18 | 18 | 126 | 136 | 175 | 26 | 166 | 255 | |
| 247 | 127 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 30 | 36 | 94 | 154 | |
| 170 | 253 | 253 | 253 | 253 | 253 | 225 | 172 | 253 | 242 | 195 | 64 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 49 | 238 | 253 | 253 | 253 | 253 | 253 | 253 | 253 | 253 | 251 | 93 | 82 | |
| 82 | 56 | 39 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 18 | 219 | 253 | |
| 253 | 253 | 253 | 253 | 198 | 182 | 247 | 241 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 80 | 156 | 107 | 253 | 253 | 205 | 11 | 0 | 43 | 154 | | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 14 | 1 | 154 | 253 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 139 | 253 | 190 | 2 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 11 | 190 | 253 | 70 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 35 | 241 | |
| 225 | 160 | 108 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 81 | 240 | 253 | 253 | 119 | 25 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 45 | 186 | 253 | 253 | 150 | 27 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 16 | 93 | 252 | 253 | 187 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 249 | 253 | 249 | 64 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 46 | 130 | 183 | 253 | |
| 253 | 207 | 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | 39 | 148 | 229 | 253 | 253 | 253 | 250 | 182 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | 0 | 0 | 0 | 0 | | | | | | | | | | | | | | | |

In [88]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - Xmin)/(Xmax-Xmin) = X/255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [89]:

```
# example data point after normalizing
print(X_train[0])
```

[illegible]

[illegible]

[illegible]

In [90]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

In [91]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:
```

```

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [92]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

In [93]:

```

# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

```

```
model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

In [94]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer ,
https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize.,
https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 111us/step - loss: 1.2659 - acc: 0.7094 - val_loss: 0.8067 - val_acc: 0.8280
Epoch 2/20
60000/60000 [=====] - 3s 42us/step - loss: 0.7120 - acc: 0.8417 - val_loss: 0.6061 - val_acc: 0.8616
Epoch 3/20
60000/60000 [=====] - 3s 43us/step - loss: 0.5854 - acc: 0.8597 - val_loss: 0.5252 - val_acc: 0.8740
Epoch 4/20
60000/60000 [=====] - 3s 43us/step - loss: 0.5244 - acc: 0.8694 - val_loss: 0.4799 - val_acc: 0.8815
Epoch 5/20
60000/60000 [=====] - 3s 42us/step - loss: 0.4871 - acc: 0.8753 - val_loss: 0.4507 - val_acc: 0.8866
Epoch 6/20
60000/60000 [=====] - 3s 43us/step - loss: 0.4614 - acc: 0.8796 - val_loss: 0.4292 - val_acc: 0.8889
Epoch 7/20
60000/60000 [=====] - 3s 42us/step - loss: 0.4423 - acc: 0.8833 - val_loss: 0.4131 - val_acc: 0.8920
Epoch 8/20
60000/60000 [=====] - 3s 44us/step - loss: 0.4274 - acc: 0.8864 - val_loss: 0.4002 - val_acc: 0.8949
Epoch 9/20
60000/60000 [=====] - 3s 43us/step - loss: 0.4154 - acc: 0.8886 - val_loss: 0.3901 - val_acc: 0.8962
Epoch 10/20
```

```

Epoch 10/20
60000/60000 [=====] - 3s 43us/step - loss: 0.4054 - acc: 0.8904 -
val_loss: 0.3811 - val_acc: 0.8976
Epoch 11/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3969 - acc: 0.8922 -
val_loss: 0.3739 - val_acc: 0.9004
Epoch 12/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3896 - acc: 0.8939 -
val_loss: 0.3671 - val_acc: 0.9017
Epoch 13/20
60000/60000 [=====] - 3s 43us/step - loss: 0.3833 - acc: 0.8952 -
val_loss: 0.3617 - val_acc: 0.9027
Epoch 14/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3777 - acc: 0.8965 -
val_loss: 0.3568 - val_acc: 0.9033
Epoch 15/20
60000/60000 [=====] - 3s 43us/step - loss: 0.3726 - acc: 0.8975 -
val_loss: 0.3524 - val_acc: 0.9046
Epoch 16/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3681 - acc: 0.8986 -
val_loss: 0.3481 - val_acc: 0.9047
Epoch 17/20
60000/60000 [=====] - 2s 41us/step - loss: 0.3640 - acc: 0.8994 -
val_loss: 0.3447 - val_acc: 0.9052
Epoch 18/20
60000/60000 [=====] - 3s 43us/step - loss: 0.3603 - acc: 0.9003 -
val_loss: 0.3414 - val_acc: 0.9069
Epoch 19/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3568 - acc: 0.9013 -
val_loss: 0.3383 - val_acc: 0.9076
Epoch 20/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3537 - acc: 0.9020 -
val_loss: 0.3357 - val_acc: 0.9086

```

In [95]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.3357250750780106

Test accuracy: 0.9086





MLP + Sigmoid activation + SGDOptimizer

In [96]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

| Layer (type) | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense_103 (Dense) | (None, 512) | 401920 |
| dense_104 (Dense) | (None, 128) | 65664 |
| dense_105 (Dense) | (None, 10) | 1290 |
| Total params: 468,874 | | |
| Trainable params: 468,874 | | |
| Non-trainable params: 0 | | |

In [97]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 116us/step - loss: 2.2665 - acc: 0.2311 - val_loss: 2.2210 - val_acc: 0.4829
Epoch 2/20
60000/60000 [=====] - 3s 46us/step - loss: 2.1767 - acc: 0.4668 - val_loss: 2.1211 - val_acc: 0.5636
Epoch 3/20
60000/60000 [=====] - 3s 47us/step - loss: 2.0609 - acc: 0.5972 - val_loss: 1.9797 - val_acc: 0.6697
Epoch 4/20
60000/60000 [=====] - 3s 46us/step - loss: 1.8947 - acc: 0.6608 - val_loss: 1.7812 - val_acc: 0.6553
Epoch 5/20
60000/60000 [=====] - 3s 47us/step - loss: 1.6747 - acc: 0.6968 - val_loss: 1.5382 - val_acc: 0.7561
Epoch 6/20
60000/60000 [=====] - 3s 47us/step - loss: 1.4344 - acc: 0.7361 - val_loss: 1.3019 - val_acc: 0.7611
Epoch 7/20
60000/60000 [=====] - 3s 46us/step - loss: 1.2204 - acc: 0.7636 - val_loss: 1.1097 - val_acc: 0.7879
Epoch 8/20
60000/60000 [=====] - 3s 46us/step - loss: 1.0534 - acc: 0.7845 - val_loss: 0.9659 - val_acc: 0.7938
Epoch 9/20
60000/60000 [=====] - 3s 46us/step - loss: 0.9287 - acc: 0.7973 - val_loss: 0.8579 - val_acc: 0.8082
Epoch 10/20
60000/60000 [=====] - 3s 45us/step - loss: 0.8215 - acc: 0.8015 - val_loss: 0.8005 - val_acc: 0.8115
```



```

60000/60000 [=====] - 3s 45us/step - loss: 0.8347 - acc: 0.8097 -
val_loss: 0.7763 - val_acc: 0.8234
Epoch 11/20
60000/60000 [=====] - 3s 46us/step - loss: 0.7623 - acc: 0.8208 -
val_loss: 0.7125 - val_acc: 0.8313
Epoch 12/20
60000/60000 [=====] - 3s 46us/step - loss: 0.7049 - acc: 0.8292 -
val_loss: 0.6622 - val_acc: 0.8421
Epoch 13/20
60000/60000 [=====] - 3s 45us/step - loss: 0.6585 - acc: 0.8369 -
val_loss: 0.6196 - val_acc: 0.8463
Epoch 14/20
60000/60000 [=====] - 3s 46us/step - loss: 0.6206 - acc: 0.8437 -
val_loss: 0.5854 - val_acc: 0.8509
Epoch 15/20
60000/60000 [=====] - 3s 46us/step - loss: 0.5888 - acc: 0.8493 -
val_loss: 0.5572 - val_acc: 0.8578
Epoch 16/20
60000/60000 [=====] - 3s 45us/step - loss: 0.5619 - acc: 0.8549 -
val_loss: 0.5329 - val_acc: 0.8620
Epoch 17/20
60000/60000 [=====] - 3s 45us/step - loss: 0.5388 - acc: 0.8593 -
val_loss: 0.5110 - val_acc: 0.8666
Epoch 18/20
60000/60000 [=====] - 3s 46us/step - loss: 0.5187 - acc: 0.8634 -
val_loss: 0.4927 - val_acc: 0.8702
Epoch 19/20
60000/60000 [=====] - 3s 46us/step - loss: 0.5012 - acc: 0.8671 -
val_loss: 0.4761 - val_acc: 0.8727
Epoch 20/20
60000/60000 [=====] - 3s 46us/step - loss: 0.4856 - acc: 0.8700 -
val_loss: 0.4625 - val_acc: 0.8764

```

In [98]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lilation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.4624546233654022

Test accuracy: 0.8764





In [99]:


```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Sigmoid activation + ADAM

In [100]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type) | Output Shape | Param # |
|--------------|--------------|---------|
|--------------|--------------|---------|

```

=====
dense_106 (Dense)                (None, 512)                401920
dense_107 (Dense)                (None, 128)                65664
dense_108 (Dense)                (None, 10)                 1290
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 125us/step - loss: 0.5394 - acc: 0.8597 -
val_loss: 0.2513 - val_acc: 0.9265
Epoch 2/20
60000/60000 [=====] - 3s 53us/step - loss: 0.2212 - acc: 0.9357 -
val_loss: 0.1897 - val_acc: 0.9416
Epoch 3/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1631 - acc: 0.9518 -
val_loss: 0.1475 - val_acc: 0.9547
Epoch 4/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1251 - acc: 0.9627 -
val_loss: 0.1142 - val_acc: 0.9649
Epoch 5/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0985 - acc: 0.9711 -
val_loss: 0.0977 - val_acc: 0.9703
Epoch 6/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0793 - acc: 0.9762 -
val_loss: 0.0865 - val_acc: 0.9729
Epoch 7/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0631 - acc: 0.9812 -
val_loss: 0.0771 - val_acc: 0.9760
Epoch 8/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0514 - acc: 0.9846 -
val_loss: 0.0805 - val_acc: 0.9754
Epoch 9/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0421 - acc: 0.9879 -
val_loss: 0.0718 - val_acc: 0.9781
Epoch 10/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0337 - acc: 0.9906 -
val_loss: 0.0650 - val_acc: 0.9789
Epoch 11/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0268 - acc: 0.9929 -
val_loss: 0.0655 - val_acc: 0.9803
Epoch 12/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0223 - acc: 0.9940 -
val_loss: 0.0668 - val_acc: 0.9797
Epoch 13/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0175 - acc: 0.9956 -
val_loss: 0.0628 - val_acc: 0.9813
Epoch 14/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0141 - acc: 0.9966 -
val_loss: 0.0650 - val_acc: 0.9818
Epoch 15/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0119 - acc: 0.9971 -
val_loss: 0.0682 - val_acc: 0.9802
Epoch 16/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0084 - acc: 0.9983 -
val_loss: 0.0679 - val_acc: 0.9812
Epoch 17/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0075 - acc: 0.9984 -
val_loss: 0.0651 - val_acc: 0.9819
Epoch 18/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0056 - acc: 0.9990 -
val_loss: 0.0647 - val_acc: 0.9823
Epoch 19/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0053 - acc: 0.9988 -
val_loss: 0.0727 - val_acc: 0.9809
Epoch 20/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0034 - acc: 0.9994 -
val_loss: 0.0700 - val_acc: 0.9826

```

In [101]:

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
```

```

print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0699703871981983
Test accuracy: 0.9826



In [102]:

```

w_after = model_sigmoid.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + ReLU +SGD

In [103]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i)}$ .
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0,\sigma) = N(0,0.062)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0,\sigma) = N(0,0.125)$ 
# out =>  $\sigma=\sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

| Layer (type) | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense_109 (Dense) | (None, 512) | 401920 |
| dense_110 (Dense) | (None, 128) | 65664 |
| dense_111 (Dense) | (None, 10) | 1290 |
| Total params: 468,874 | | |
| Trainable params: 468,874 | | |
| Non-trainable params: 0 | | |

In [104]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 120us/step - loss: 0.7727 - acc: 0.7796 - val_loss: 0.3984 - val_acc: 0.8890
Epoch 2/20
60000/60000 [=====] - 3s 47us/step - loss: 0.3619 - acc: 0.8984 - val_loss: 0.3029 - val_acc: 0.9131
Epoch 3/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2946 - acc: 0.9168 - val_loss: 0.2617 - val_acc: 0.9263
```

```

Epoch 4/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2585 - acc: 0.9263 -
val_loss: 0.2349 - val_acc: 0.9343
Epoch 5/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2330 - acc: 0.9342 -
val_loss: 0.2166 - val_acc: 0.9387
Epoch 6/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2137 - acc: 0.9399 -
val_loss: 0.2030 - val_acc: 0.9409
Epoch 7/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1983 - acc: 0.9437 -
val_loss: 0.1905 - val_acc: 0.9453
Epoch 8/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1855 - acc: 0.9478 -
val_loss: 0.1810 - val_acc: 0.9477
Epoch 9/20
60000/60000 [=====] - 3s 46us/step - loss: 0.1747 - acc: 0.9501 -
val_loss: 0.1721 - val_acc: 0.9485
Epoch 10/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1651 - acc: 0.9534 -
val_loss: 0.1646 - val_acc: 0.9510
Epoch 11/20
60000/60000 [=====] - 3s 45us/step - loss: 0.1567 - acc: 0.9556 -
val_loss: 0.1579 - val_acc: 0.9527
Epoch 12/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1494 - acc: 0.9573 -
val_loss: 0.1536 - val_acc: 0.9545
Epoch 13/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1424 - acc: 0.9602 -
val_loss: 0.1474 - val_acc: 0.9564
Epoch 14/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1366 - acc: 0.9618 -
val_loss: 0.1424 - val_acc: 0.9571
Epoch 15/20
60000/60000 [=====] - 3s 49us/step - loss: 0.1311 - acc: 0.9636 -
val_loss: 0.1413 - val_acc: 0.9574
Epoch 16/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1259 - acc: 0.9649 -
val_loss: 0.1350 - val_acc: 0.9590
Epoch 17/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1211 - acc: 0.9660 -
val_loss: 0.1317 - val_acc: 0.9595
Epoch 18/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1166 - acc: 0.9677 -
val_loss: 0.1276 - val_acc: 0.9609
Epoch 19/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1125 - acc: 0.9689 -
val_loss: 0.1260 - val_acc: 0.9610
Epoch 20/20
60000/60000 [=====] - 3s 46us/step - loss: 0.1087 - acc: 0.9698 -
val_loss: 0.1225 - val_acc: 0.9627

```

In [105]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lida
tion_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12253610298782587
Test accuracy: 0.9627



In [106]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU + ADAM

In [107]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type) | Output Shape | Param # |
|-------------------|--------------|---------|
| dense_112 (Dense) | (None, 512) | 401920 |
| dense_113 (Dense) | (None, 128) | 65664 |
| dense_114 (Dense) | (None, 10) | 1290 |

=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
=====

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 130us/step - loss: 0.2288 - acc: 0.9326 - val_loss: 0.1132 - val_acc: 0.9655

Epoch 2/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0854 - acc: 0.9737 - val_loss: 0.0770 - val_acc: 0.9761

Epoch 3/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0522 - acc: 0.9837 - val_loss: 0.0829 - val_acc: 0.9741

Epoch 4/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0380 - acc: 0.9879 - val_loss: 0.0755 - val_acc: 0.9764

Epoch 5/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0262 - acc: 0.9918 - val_loss: 0.0712 - val_acc: 0.9793

Epoch 6/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0200 - acc: 0.9940 - val_loss: 0.0768 - val_acc: 0.9788

Epoch 7/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0152 - acc: 0.9951 - val_loss: 0.0642 - val_acc: 0.9828

Epoch 8/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0151 - acc: 0.9950 - val_loss: 0.0781 - val_acc: 0.9796

Epoch 9/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0151 - acc: 0.9950 - val_loss: 0.0795 - val_acc: 0.9795

Epoch 10/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0111 - acc: 0.9962 - val_loss: 0.0853 - val_acc: 0.9806

Epoch 11/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0105 - acc: 0.9965 - val_loss: 0.1058 - val_acc: 0.9755

Epoch 12/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0129 - acc: 0.9956 - val_loss: 0.0834 - val_acc: 0.9799

Epoch 13/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0072 - acc: 0.9977 - val_loss: 0.0840 - val_acc: 0.9814

Epoch 14/20

60000/60000 [=====] - 3s 53us/step - loss: 0.0095 - acc: 0.9968 - val_loss: 0.0819 - val_acc: 0.9814

Epoch 15/20


```

60000/60000 [=====] - 3s 52us/step - loss: 0.0096 - acc: 0.9967 -
val_loss: 0.1173 - val_acc: 0.9756
Epoch 16/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0096 - acc: 0.9966 -
val_loss: 0.0950 - val_acc: 0.9800
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0068 - acc: 0.9978 -
val_loss: 0.1090 - val_acc: 0.9777
Epoch 18/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0093 - acc: 0.9972 -
val_loss: 0.1086 - val_acc: 0.9780
Epoch 19/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0082 - acc: 0.9971 -
val_loss: 0.0813 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0038 - acc: 0.9988 -
val_loss: 0.0990 - val_acc: 0.9826

```

In [108]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09899708161438571

Test accuracy: 0.9826



In [109]:

```
w_after = model_relu.get_weights()
```

```


n1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [110]:

```

# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.055, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_115 (Dense) | (None, 512) | 401920 |
| batch_normalization_45 (Batch Normalization) | (None, 512) | 2048 |
| dense_116 (Dense) | (None, 128) | 65664 |
| batch_normalization_46 (Batch Normalization) | (None, 128) | 512 |
| dense_117 (Dense) | (None, 10) | 1290 |
| Total params: 471,434 | | |
| Trainable params: 470,154 | | |
| Non-trainable params: 1,280 | | |

In [111]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 10s 163us/step - loss: 0.3013 - acc: 0.9113 - val_loss: 0.2207 - val_acc: 0.9347
Epoch 2/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1757 - acc: 0.9487 - val_loss: 0.1698 - val_acc: 0.9502
Epoch 3/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1382 - acc: 0.9593 - val_loss: 0.1542 - val_acc: 0.9546
Epoch 4/20
60000/60000 [=====] - 5s 80us/step - loss: 0.1140 - acc: 0.9656 - val_loss: 0.1405 - val_acc: 0.9581
Epoch 5/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0947 - acc: 0.9720 - val_loss: 0.1276 - val_acc: 0.9625
Epoch 6/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0804 - acc: 0.9753 - val_loss: 0.1157 - val_acc: 0.9654
Epoch 7/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0674 - acc: 0.9795 - val_loss: 0.1131 - val_acc: 0.9656
Epoch 8/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0582 - acc: 0.9826 - val_loss: 0.1060 - val_acc: 0.9675
Epoch 9/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0509 - acc: 0.9842 - val_loss: 0.1078 - val_acc: 0.9687
Epoch 10/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0474 - acc: 0.9853 - val_loss: 0.1032 - val_acc: 0.9681
Epoch 11/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0395 - acc: 0.9876 - val_loss: 0.0999 - val_acc: 0.9705
Epoch 12/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0341 - acc: 0.9889 - val_loss: 0.0963 - val_acc: 0.9719
Epoch 13/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0308 - acc: 0.9902 - val_loss: 0.0936 - val_acc: 0.9710
Epoch 14/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0265 - acc: 0.9917 - val_loss: 0.0916 - val_acc: 0.9733
Epoch 15/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0224 - acc: 0.9933 - val_loss: 0.1015 - val_acc: 0.9717
Epoch 16/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0202 - acc: 0.9936 - val_loss: 0.0969 - val_acc: 0.9737
Epoch 17/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0208 - acc: 0.9932 - val_loss: 0.1138 - val_acc: 0.9698
```

```

val_loss: 0.0960 - val_acc: 0.9750
Epoch 18/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0196 - acc: 0.9936 -
val_loss: 0.0960 - val_acc: 0.9750
Epoch 19/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0177 - acc: 0.9941 -
val_loss: 0.0924 - val_acc: 0.9749
Epoch 20/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0169 - acc: 0.9944 -
val_loss: 0.0983 - val_acc: 0.9734

```

In [112]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09825632937119808
Test accuracy: 0.9734



In [113]:

```

w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")

```


```

ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



5. MLP + Dropout + AdamOptimizer

In [114]:

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_118 (Dense) | (None, 512) | 401920 |
| batch_normalization_47 (Batch Normalization) | (None, 512) | 2048 |
| dropout_35 (Dropout) | (None, 512) | 0 |
| dense_119 (Dense) | (None, 128) | 65664 |
| batch_normalization_48 (Batch Normalization) | (None, 128) | 512 |
| dropout_36 (Dropout) | (None, 128) | 0 |

| | | |
|-------------------|------------|------|
| dense_120 (Dense) | (None, 10) | 1290 |
|-------------------|------------|------|

=====
Total params: 471,434
Trainable params: 470,154
Non-trainable params: 1,280
=====

In [115]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 177us/step - loss: 0.6745 - acc: 0.7926 - val_loss: 0.2803 - val_acc: 0.9194
Epoch 2/20
60000/60000 [=====] - 5s 85us/step - loss: 0.4344 - acc: 0.8685 - val_loss: 0.2526 - val_acc: 0.9237
Epoch 3/20
60000/60000 [=====] - 5s 84us/step - loss: 0.3827 - acc: 0.8839 - val_loss: 0.2336 - val_acc: 0.9312
Epoch 4/20
60000/60000 [=====] - 6s 98us/step - loss: 0.3548 - acc: 0.8927 - val_loss: 0.2309 - val_acc: 0.9323
Epoch 5/20
60000/60000 [=====] - 5s 86us/step - loss: 0.3339 - acc: 0.8993 - val_loss: 0.2153 - val_acc: 0.9359
Epoch 6/20
60000/60000 [=====] - 5s 84us/step - loss: 0.3227 - acc: 0.9015 - val_loss: 0.2070 - val_acc: 0.9396
Epoch 7/20
60000/60000 [=====] - 5s 85us/step - loss: 0.3092 - acc: 0.9063 - val_loss: 0.1990 - val_acc: 0.9406
Epoch 8/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2952 - acc: 0.9118 - val_loss: 0.1854 - val_acc: 0.9429
Epoch 9/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2866 - acc: 0.9136 - val_loss: 0.1767 - val_acc: 0.9462
Epoch 10/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2704 - acc: 0.9189 - val_loss: 0.1742 - val_acc: 0.9467
Epoch 11/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2585 - acc: 0.9223 - val_loss: 0.1670 - val_acc: 0.9505
Epoch 12/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2499 - acc: 0.9249 - val_loss: 0.1550 - val_acc: 0.9541
Epoch 13/20
60000/60000 [=====] - 5s 84us/step - loss: 0.2384 - acc: 0.9281 - val_loss: 0.1533 - val_acc: 0.9534
Epoch 14/20
60000/60000 [=====] - 5s 84us/step - loss: 0.2262 - acc: 0.9318 - val_loss: 0.1447 - val_acc: 0.9560
Epoch 15/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2201 - acc: 0.9342 - val_loss: 0.1387 - val_acc: 0.9570
Epoch 16/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2133 - acc: 0.9354 - val_loss: 0.1286 - val_acc: 0.9619
Epoch 17/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1993 - acc: 0.9406 - val_loss: 0.1242 - val_acc: 0.9622
Epoch 18/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1947 - acc: 0.9415 - val_loss: 0.1251 - val_acc: 0.9626
Epoch 19/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1854 - acc: 0.9446 - val_loss: 0.1217 - val_acc: 0.9649
Epoch 20/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1799 - acc: 0.9449 - val_loss: 0.1129 - val_acc: 0.9685
```

In [116]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.11289959549605846

Test accuracy: 0.9685



In [117]:

```
w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
```

```
plt.figure(figsize=(10, 10))
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hyper-parameter tuning of Keras models using Sklearn

In [118]:

```
from keras.optimizers import Adam, RMSprop, SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

In [119]:

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid', 'relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)

/opt/conda/lib/python3.6/site-packages/sklearn/model_selection/_split.py:1978: FutureWarning: The default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this warning.
  warnings.warn(CV_WARNING, FutureWarning)
```

In [120]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
```



```

means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

```

```

Best: 0.976300 using {'activ': 'relu'}
0.975617 (0.001652) with: {'activ': 'sigmoid'}
0.976300 (0.002729) with: {'activ': 'relu'}

```

6. Model 1 with two hidden layers + Batch Normalization + Dropout:

In [121]:

```

model_relu = Sequential()
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_142 (Dense) | (None, 256) | 200960 |
| batch_normalization_49 (Batch Normalization) | (None, 256) | 1024 |
| dropout_37 (Dropout) | (None, 256) | 0 |
| dense_143 (Dense) | (None, 128) | 32896 |
| batch_normalization_50 (Batch Normalization) | (None, 128) | 512 |
| dropout_38 (Dropout) | (None, 128) | 0 |
| dense_144 (Dense) | (None, 10) | 1290 |

Total params: 236,682

Trainable params: 235,914

Non-trainable params: 768

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 192us/step - loss: 0.3786 - acc: 0.8847 - val_loss: 0.1445 - val_acc: 0.9560

Epoch 2/20

60000/60000 [=====] - 5s 88us/step - loss: 0.1768 - acc: 0.9470 - val_loss: 0.1059 - val_acc: 0.9682

Epoch 3/20

60000/60000 [=====] - 5s 89us/step - loss: 0.1326 - acc: 0.9583 - val_loss: 0.0858 - val_acc: 0.9728

Epoch 4/20

60000/60000 [=====] - 5s 88us/step - loss: 0.1104 - acc: 0.9652 - val_loss: 0.0835 - val_acc: 0.9741

Epoch 5/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0954 - acc: 0.9700 - val_loss: 0.0715 - val_acc: 0.9782

Epoch 6/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0858 - acc: 0.9732 - val_loss: 0.0671 - val_acc: 0.9801

Epoch 7/20

```

60000/60000 [=====] - 5s 88us/step - loss: 0.0755 - acc: 0.9759 -
val_loss: 0.0704 - val_acc: 0.9786
Epoch 8/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0683 - acc: 0.9783 -
val_loss: 0.0650 - val_acc: 0.9805
Epoch 9/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0652 - acc: 0.9794 -
val_loss: 0.0627 - val_acc: 0.9810
Epoch 10/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0610 - acc: 0.9803 -
val_loss: 0.0710 - val_acc: 0.9803
Epoch 11/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0559 - acc: 0.9811 -
val_loss: 0.0630 - val_acc: 0.9816
Epoch 12/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0525 - acc: 0.9830 -
val_loss: 0.0640 - val_acc: 0.9814
Epoch 13/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0500 - acc: 0.9840 -
val_loss: 0.0710 - val_acc: 0.9798
Epoch 14/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0459 - acc: 0.9850 -
val_loss: 0.0672 - val_acc: 0.9807
Epoch 15/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0456 - acc: 0.9844 -
val_loss: 0.0621 - val_acc: 0.9820
Epoch 16/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0428 - acc: 0.9859 -
val_loss: 0.0575 - val_acc: 0.9826
Epoch 17/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0427 - acc: 0.9861 -
val_loss: 0.0585 - val_acc: 0.9826
Epoch 18/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0401 - acc: 0.9862 -
val_loss: 0.0611 - val_acc: 0.9833
Epoch 19/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0387 - acc: 0.9876 -
val_loss: 0.0600 - val_acc: 0.9818
Epoch 20/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0358 - acc: 0.9880 -
val_loss: 0.0637 - val_acc: 0.9827

```

In [122]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06367886031113303

Test accuracy: 0.9827



In [123]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



7. Model 2 with three hidden layers + Batch Normalization + Dropout:

In [124]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
```

```

model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| ===== | ===== | ===== |
| dense_145 (Dense) | (None, 512) | 401920 |
| batch_normalization_51 (Batch Normalization) | (None, 512) | 2048 |
| dropout_39 (Dropout) | (None, 512) | 0 |
| dense_146 (Dense) | (None, 256) | 131328 |
| batch_normalization_52 (Batch Normalization) | (None, 256) | 1024 |
| dropout_40 (Dropout) | (None, 256) | 0 |
| dense_147 (Dense) | (None, 128) | 32896 |
| batch_normalization_53 (Batch Normalization) | (None, 128) | 512 |
| dropout_41 (Dropout) | (None, 128) | 0 |
| dense_148 (Dense) | (None, 10) | 1290 |
| ===== | ===== | ===== |
| Total params: 571,018 | | |
| Trainable params: 569,226 | | |
| Non-trainable params: 1,792 | | |

```

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 13s 224us/step - loss: 0.3627 - acc: 0.8891 - val_loss: 0.1164 - val_acc: 0.9641
Epoch 2/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1567 - acc: 0.9530 - val_loss: 0.0919 - val_acc: 0.9695
Epoch 3/20
60000/60000 [=====] - 6s 108us/step - loss: 0.1184 - acc: 0.9636 - val_loss: 0.0804 - val_acc: 0.9751
Epoch 4/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1020 - acc: 0.9687 - val_loss: 0.0829 - val_acc: 0.9746
Epoch 5/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0844 - acc: 0.9733 - val_loss: 0.0787 - val_acc: 0.9756
Epoch 6/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0770 - acc: 0.9754 - val_loss: 0.0699 - val_acc: 0.9785
Epoch 7/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0689 - acc: 0.9780 - val_loss: 0.0686 - val_acc: 0.9770
Epoch 8/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0626 - acc: 0.9799 - val_loss: 0.0690 - val_acc: 0.9794
Epoch 9/20
60000/60000 [=====] - 7s 111us/step - loss: 0.0568 - acc: 0.9819 - val_loss: 0.0632 - val_acc: 0.9808
Epoch 10/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0511 - acc: 0.9834 - val_loss: 0.0666 - val_acc: 0.9802

```

```

Epoch 11/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0495 - acc: 0.9839 -
val_loss: 0.0640 - val_acc: 0.9800
Epoch 12/20
60000/60000 [=====] - 7s 109us/step - loss: 0.0475 - acc: 0.9845 -
val_loss: 0.0567 - val_acc: 0.9833
Epoch 13/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0444 - acc: 0.9856 -
val_loss: 0.0566 - val_acc: 0.9836
Epoch 14/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0388 - acc: 0.9875 -
val_loss: 0.0595 - val_acc: 0.9823
Epoch 15/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0391 - acc: 0.9874 -
val_loss: 0.0619 - val_acc: 0.9820
Epoch 16/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0378 - acc: 0.9879 -
val_loss: 0.0609 - val_acc: 0.9828
Epoch 17/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0332 - acc: 0.9890 -
val_loss: 0.0576 - val_acc: 0.9844
Epoch 18/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0349 - acc: 0.9884 -
val_loss: 0.0621 - val_acc: 0.9821
Epoch 19/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0323 - acc: 0.9894 -
val_loss: 0.0549 - val_acc: 0.9857
Epoch 20/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0317 - acc: 0.9896 -
val_loss: 0.0609 - val_acc: 0.9843

```

In [125]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.060851101654174275
Test accuracy: 0.9843





In [126]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



8. Model 3 with five hidden layers + Batch Normalization + Dropout:

In [127]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
```

```

model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125,
seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_149 (Dense) | (None, 512) | 401920 |
| batch_normalization_54 (Batch Normalization) | (None, 512) | 2048 |
| dropout_42 (Dropout) | (None, 512) | 0 |
| dense_150 (Dense) | (None, 256) | 131328 |
| batch_normalization_55 (Batch Normalization) | (None, 256) | 1024 |
| dropout_43 (Dropout) | (None, 256) | 0 |
| dense_151 (Dense) | (None, 128) | 32896 |
| batch_normalization_56 (Batch Normalization) | (None, 128) | 512 |
| dropout_44 (Dropout) | (None, 128) | 0 |
| dense_152 (Dense) | (None, 64) | 8256 |
| batch_normalization_57 (Batch Normalization) | (None, 64) | 256 |
| dense_153 (Dense) | (None, 32) | 2080 |
| batch_normalization_58 (Batch Normalization) | (None, 32) | 128 |
| dense_154 (Dense) | (None, 10) | 330 |
| Total params: 580,778 | | |
| Trainable params: 578,794 | | |
| Non-trainable params: 1,984 | | |

```

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 16s 273us/step - loss: 0.4450 - acc: 0.8676 - val_loss: 0.1393 - val_acc: 0.9595
Epoch 2/20
60000/60000 [=====] - 9s 144us/step - loss: 0.1794 - acc: 0.9467 - val_loss: 0.1054 - val_acc: 0.9682
Epoch 3/20
60000/60000 [=====] - 9s 143us/step - loss: 0.1373 - acc: 0.9597 - val_loss: 0.0938 - val_acc: 0.9715
Epoch 4/20
60000/60000 [=====] - 9s 144us/step - loss: 0.1131 - acc: 0.9658 - val_loss: 0.0848 - val_acc: 0.9747
Epoch 5/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0990 - acc: 0.9699 - val_loss: 0.0750 - val_acc: 0.9779
Epoch 6/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0866 - acc: 0.9733 - val_loss: 0.0688 - val_acc: 0.9805
Epoch 7/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0795 - acc: 0.9754 - val_loss: 0.0703 - val_acc: 0.9788
Epoch 8/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0718 - acc: 0.9781 - val_loss: 0.0624 - val_acc: 0.9816
Epoch 9/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0664 - acc: 0.9800 - val_loss: 0.0672 - val_acc: 0.9805
Epoch 10/20

```

```

Epoch 10/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0625 - acc: 0.9800 -
val_loss: 0.0611 - val_acc: 0.9813
Epoch 11/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0588 - acc: 0.9817 -
val_loss: 0.0627 - val_acc: 0.9824
Epoch 12/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0559 - acc: 0.9827 -
val_loss: 0.0598 - val_acc: 0.9830
Epoch 13/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0528 - acc: 0.9834 -
val_loss: 0.0611 - val_acc: 0.9839
Epoch 14/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0510 - acc: 0.9841 -
val_loss: 0.0659 - val_acc: 0.9822
Epoch 15/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0460 - acc: 0.9860 -
val_loss: 0.0639 - val_acc: 0.9822
Epoch 16/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0444 - acc: 0.9861 -
val_loss: 0.0593 - val_acc: 0.9834
Epoch 17/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0433 - acc: 0.9865 -
val_loss: 0.0599 - val_acc: 0.9835
Epoch 18/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0391 - acc: 0.9880 -
val_loss: 0.0582 - val_acc: 0.9850
Epoch 19/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0407 - acc: 0.9875 -
val_loss: 0.0593 - val_acc: 0.9835
Epoch 20/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0376 - acc: 0.9880 -
val_loss: 0.0561 - val_acc: 0.9843

```

In [128]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05613952756321523

Test accuracy: 0.9843





In [129]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Conclusion:

In [150]:

```
from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["# Layers", "Epoch", "Accuracy"]
x.add_row(["2", 20, 0.9827])
x.add_row(["3", 20, 0.9843])
x.add_row(["5", 20, 0.9843])
print(x)
```

```
+-----+-----+-----+
| # Layers | Epoch | Accuracy |
+-----+-----+-----+
|      2   |    20 |   0.9827  |
```

| | | |
|---|----|--------|
| 3 | 20 | 0.9843 |
| 5 | 20 | 0.9843 |

9. Model 1: MLP + BatchNormalization + Dropout (0.30)

- #layers: 5
- activation: ReLU
- Weight Initializer: RandomNormal
- Optimizer: ADAM

In [131]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.050, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_155 (Dense) | (None, 512) | 401920 |
| batch_normalization_59 (Batch Normalization) | (None, 512) | 2048 |
| dropout_45 (Dropout) | (None, 512) | 0 |
| dense_156 (Dense) | (None, 256) | 131328 |
| batch_normalization_60 (Batch Normalization) | (None, 256) | 1024 |
| dropout_46 (Dropout) | (None, 256) | 0 |
| dense_157 (Dense) | (None, 128) | 32896 |
| batch_normalization_61 (Batch Normalization) | (None, 128) | 512 |
| dropout_47 (Dropout) | (None, 128) | 0 |
| dense_158 (Dense) | (None, 64) | 8256 |
| batch_normalization_62 (Batch Normalization) | (None, 64) | 256 |
| dense_159 (Dense) | (None, 32) | 2080 |
| batch_normalization_63 (Batch Normalization) | (None, 32) | 128 |
| dense_160 (Dense) | (None, 10) | 330 |
| Total params: 580,778 | | |

Trainable params: 518,194
Non-trainable params: 1,984

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 278us/step - loss: 0.4477 - acc: 0.8703 - val_loss: 0.1462 - val_acc: 0.9558

Epoch 2/20

60000/60000 [=====] - 8s 142us/step - loss: 0.1781 - acc: 0.9478 - val_loss: 0.0991 - val_acc: 0.9717

Epoch 3/20

60000/60000 [=====] - 9s 158us/step - loss: 0.1386 - acc: 0.9588 - val_loss: 0.0854 - val_acc: 0.9741

Epoch 4/20

60000/60000 [=====] - 9s 144us/step - loss: 0.1147 - acc: 0.9652 - val_loss: 0.0748 - val_acc: 0.9775

Epoch 5/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0966 - acc: 0.9700 - val_loss: 0.0822 - val_acc: 0.9767

Epoch 6/20

60000/60000 [=====] - 8s 141us/step - loss: 0.0863 - acc: 0.9734 - val_loss: 0.0783 - val_acc: 0.9772

Epoch 7/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0803 - acc: 0.9758 - val_loss: 0.0704 - val_acc: 0.9790

Epoch 8/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0723 - acc: 0.9777 - val_loss: 0.0646 - val_acc: 0.9802

Epoch 9/20

60000/60000 [=====] - 8s 142us/step - loss: 0.0678 - acc: 0.9787 - val_loss: 0.0602 - val_acc: 0.9831

Epoch 10/20

60000/60000 [=====] - 8s 142us/step - loss: 0.0644 - acc: 0.9802 - val_loss: 0.0638 - val_acc: 0.9814

Epoch 11/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0583 - acc: 0.9819 - val_loss: 0.0568 - val_acc: 0.9834

Epoch 12/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0559 - acc: 0.9826 - val_loss: 0.0593 - val_acc: 0.9831

Epoch 13/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0529 - acc: 0.9838 - val_loss: 0.0670 - val_acc: 0.9802

Epoch 14/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0502 - acc: 0.9838 - val_loss: 0.0541 - val_acc: 0.9845

Epoch 15/20

60000/60000 [=====] - 8s 141us/step - loss: 0.0468 - acc: 0.9851 - val_loss: 0.0590 - val_acc: 0.9848

Epoch 16/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0438 - acc: 0.9863 - val_loss: 0.0633 - val_acc: 0.9835

Epoch 17/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0444 - acc: 0.9860 - val_loss: 0.0622 - val_acc: 0.9836

Epoch 18/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0423 - acc: 0.9872 - val_loss: 0.0568 - val_acc: 0.9834

Epoch 19/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0396 - acc: 0.9871 - val_loss: 0.0584 - val_acc: 0.9849

Epoch 20/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0387 - acc: 0.9882 - val_loss: 0.0600 - val_acc: 0.9832

In [132]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
```

```

x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06003061625857372

Test accuracy: 0.9832



In [133]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).

max_open_warning, RuntimeWarning)



10. Model 2: MLP + Dropout (0.30)

- #layers: 5
- activation: ReLU
- Weight Initializer: He Normal
- Optimizer: ADAM

In [134]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal()))
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='relu', kernel_initializer=he_normal()))
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', kernel_initializer=he_normal()))
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='relu', kernel_initializer=he_normal()))
model_relu.add(Dense(32, activation='relu', kernel_initializer=he_normal()))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type) | Output Shape | Param # |
|---|--------------|---------|
| ===== | | |
| dense_161 (Dense) | (None, 512) | 401920 |
| dropout_48 (Dropout) | (None, 512) | 0 |
| dense_162 (Dense) | (None, 256) | 131328 |
| dropout_49 (Dropout) | (None, 256) | 0 |
| dense_163 (Dense) | (None, 128) | 32896 |
| dropout_50 (Dropout) | (None, 128) | 0 |
| dense_164 (Dense) | (None, 64) | 8256 |
| dense_165 (Dense) | (None, 32) | 2080 |
| dense_166 (Dense) | (None, 10) | 330 |
| ===== | | |
| Total params: 576,810 | | |
| Trainable params: 576,810 | | |
| Non-trainable params: 0 | | |
| None | | |
| Train on 60000 samples, validate on 10000 samples | | |

```

train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 12s 194us/step - loss: 0.4457 - acc: 0.8628 - val_loss: 0.1451 - val_acc: 0.9572
Epoch 2/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1724 - acc: 0.9502 - val_loss: 0.1117 - val_acc: 0.9678
Epoch 3/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1319 - acc: 0.9628 - val_loss: 0.0957 - val_acc: 0.9693
Epoch 4/20
60000/60000 [=====] - 4s 75us/step - loss: 0.1059 - acc: 0.9694 - val_loss: 0.0755 - val_acc: 0.9762
Epoch 5/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0885 - acc: 0.9738 - val_loss: 0.0729 - val_acc: 0.9785
Epoch 6/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0804 - acc: 0.9756 - val_loss: 0.0725 - val_acc: 0.9783
Epoch 7/20
60000/60000 [=====] - 4s 75us/step - loss: 0.0689 - acc: 0.9796 - val_loss: 0.0730 - val_acc: 0.9782
Epoch 8/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0650 - acc: 0.9807 - val_loss: 0.0749 - val_acc: 0.9788
Epoch 9/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0606 - acc: 0.9816 - val_loss: 0.0718 - val_acc: 0.9794
Epoch 10/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0548 - acc: 0.9834 - val_loss: 0.0708 - val_acc: 0.9813
Epoch 11/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0544 - acc: 0.9832 - val_loss: 0.0816 - val_acc: 0.9791
Epoch 12/20
60000/60000 [=====] - 4s 75us/step - loss: 0.0469 - acc: 0.9856 - val_loss: 0.0659 - val_acc: 0.9825
Epoch 13/20
60000/60000 [=====] - 5s 76us/step - loss: 0.0457 - acc: 0.9866 - val_loss: 0.0730 - val_acc: 0.9813
Epoch 14/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0411 - acc: 0.9875 - val_loss: 0.0662 - val_acc: 0.9829
Epoch 15/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0410 - acc: 0.9877 - val_loss: 0.0734 - val_acc: 0.9827
Epoch 16/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0397 - acc: 0.9886 - val_loss: 0.0705 - val_acc: 0.9828
Epoch 17/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0360 - acc: 0.9892 - val_loss: 0.0742 - val_acc: 0.9814
Epoch 18/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0390 - acc: 0.9885 - val_loss: 0.0655 - val_acc: 0.9829
Epoch 19/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0309 - acc: 0.9904 - val_loss: 0.0689 - val_acc: 0.9833
Epoch 20/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0335 - acc: 0.9901 - val_loss: 0.0675 - val_acc: 0.9833

```

In [135]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va

```

```

# history = model_deep.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06747149933629462

Test accuracy: 0.9833



In [136]:

```

w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```





11. Model 3: MLP + BatchNormalization + Dropout (0.40)

- #layers: 5
- activation: ReLU
- Weight Initializer: RandomNormal
- Optimizer: ADAM

In [137]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.050, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))
model_relu.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_167 (Dense) | (None, 512) | 401920 |
| batch_normalization_64 (Batch Normalization) | (None, 512) | 2048 |
| dropout_51 (Dropout) | (None, 512) | 0 |
| dense_168 (Dense) | (None, 256) | 131328 |
| batch_normalization_65 (Batch Normalization) | (None, 256) | 1024 |
| dropout_52 (Dropout) | (None, 256) | 0 |
| dense_169 (Dense) | (None, 128) | 32896 |
| batch_normalization_66 (Batch Normalization) | (None, 128) | 512 |
| dropout_53 (Dropout) | (None, 128) | 0 |
| dense_170 (Dense) | (None, 64) | 8256 |
| batch_normalization_67 (Batch Normalization) | (None, 64) | 256 |
| dense_171 (Dense) | (None, 32) | 2080 |
| batch_normalization_68 (Batch Normalization) | (None, 32) | 128 |


```
dense_172 (Dense)                (None, 10)                330
=====
Total params: 580,778
Trainable params: 578,794
Non-trainable params: 1,984
```

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 291us/step - loss: 0.5734 - acc: 0.8261 - val_loss: 0.1564 - val_acc: 0.9542

Epoch 2/20

60000/60000 [=====] - 9s 145us/step - loss: 0.2228 - acc: 0.9344 - val_loss: 0.1220 - val_acc: 0.9634

Epoch 3/20

60000/60000 [=====] - 9s 153us/step - loss: 0.1739 - acc: 0.9496 - val_loss: 0.1009 - val_acc: 0.9690

Epoch 4/20

60000/60000 [=====] - 9s 148us/step - loss: 0.1416 - acc: 0.9580 - val_loss: 0.0903 - val_acc: 0.9737

Epoch 5/20

60000/60000 [=====] - 9s 145us/step - loss: 0.1279 - acc: 0.9622 - val_loss: 0.0767 - val_acc: 0.9766

Epoch 6/20

60000/60000 [=====] - 9s 145us/step - loss: 0.1174 - acc: 0.9654 - val_loss: 0.0806 - val_acc: 0.9762

Epoch 7/20

60000/60000 [=====] - 9s 145us/step - loss: 0.1058 - acc: 0.9683 - val_loss: 0.0778 - val_acc: 0.9760

Epoch 8/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0971 - acc: 0.9710 - val_loss: 0.0673 - val_acc: 0.9793

Epoch 9/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0893 - acc: 0.9730 - val_loss: 0.0737 - val_acc: 0.9785

Epoch 10/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0840 - acc: 0.9748 - val_loss: 0.0634 - val_acc: 0.9824

Epoch 11/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0816 - acc: 0.9757 - val_loss: 0.0611 - val_acc: 0.9824

Epoch 12/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0751 - acc: 0.9776 - val_loss: 0.0667 - val_acc: 0.9811

Epoch 13/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0715 - acc: 0.9783 - val_loss: 0.0640 - val_acc: 0.9810

Epoch 14/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0658 - acc: 0.9798 - val_loss: 0.0632 - val_acc: 0.9824

Epoch 15/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0665 - acc: 0.9799 - val_loss: 0.0613 - val_acc: 0.9835

Epoch 16/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0623 - acc: 0.9809 - val_loss: 0.0619 - val_acc: 0.9819

Epoch 17/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0597 - acc: 0.9822 - val_loss: 0.0571 - val_acc: 0.9836

Epoch 18/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0577 - acc: 0.9822 - val_loss: 0.0557 - val_acc: 0.9836

Epoch 19/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0558 - acc: 0.9826 - val_loss: 0.0590 - val_acc: 0.9833

Epoch 20/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0522 - acc: 0.9837 - val_loss: 0.0556 - val_acc: 0.9836

In [138]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.055616749329864976

Test accuracy: 0.9836

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).



In [139]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")

```

```
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



12. Model 4: MLP + BatchNormalization + Dropout (0.30)

- #layers: 5
- activation: sigmoid
- Weight Initializer: RandomNormal
- Optimizer: ADAM

In [140]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.072, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.102, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.144, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_173 (Dense) | (None, 512) | 401920 |
| batch_normalization_69 (Batch Normalization) | (None, 512) | 2048 |
| dropout_54 (Dropout) | (None, 512) | 0 |
| dense_174 (Dense) | (None, 256) | 131328 |
| batch_normalization_70 (Batch Normalization) | (None, 256) | 1024 |
| dropout_55 (Dropout) | (None, 256) | 0 |

| | | |
|--|-------------|-------|
| dropout_55 (Dropout) | (None, 256) | 0 |
| dense_175 (Dense) | (None, 128) | 32896 |
| batch_normalization_71 (Batch Normalization) | (None, 128) | 512 |
| dropout_56 (Dropout) | (None, 128) | 0 |
| dense_176 (Dense) | (None, 64) | 8256 |
| batch_normalization_72 (Batch Normalization) | (None, 64) | 256 |
| dense_177 (Dense) | (None, 32) | 2080 |
| batch_normalization_73 (Batch Normalization) | (None, 32) | 128 |
| dense_178 (Dense) | (None, 10) | 330 |
| ===== | | |
| Total params: 580,778 | | |
| Trainable params: 578,794 | | |
| Non-trainable params: 1,984 | | |

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 18s 303us/step - loss: 0.4027 - acc: 0.8826 - val_loss: 0.2217 - val_acc: 0.9337

Epoch 2/20

60000/60000 [=====] - 9s 147us/step - loss: 0.2438 - acc: 0.9262 - val_loss: 0.1567 - val_acc: 0.9532

Epoch 3/20

60000/60000 [=====] - 9s 147us/step - loss: 0.1918 - acc: 0.9430 - val_loss: 0.1241 - val_acc: 0.9624

Epoch 4/20

60000/60000 [=====] - 9s 146us/step - loss: 0.1679 - acc: 0.9493 - val_loss: 0.1181 - val_acc: 0.9643

Epoch 5/20

60000/60000 [=====] - 9s 146us/step - loss: 0.1455 - acc: 0.9564 - val_loss: 0.1064 - val_acc: 0.9688

Epoch 6/20

60000/60000 [=====] - 9s 147us/step - loss: 0.1307 - acc: 0.9600 - val_loss: 0.0982 - val_acc: 0.9722

Epoch 7/20

60000/60000 [=====] - 9s 147us/step - loss: 0.1157 - acc: 0.9651 - val_loss: 0.0866 - val_acc: 0.9719

Epoch 8/20

60000/60000 [=====] - 9s 147us/step - loss: 0.1076 - acc: 0.9669 - val_loss: 0.0824 - val_acc: 0.9745

Epoch 9/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0948 - acc: 0.9706 - val_loss: 0.0808 - val_acc: 0.9747

Epoch 10/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0923 - acc: 0.9713 - val_loss: 0.0797 - val_acc: 0.9771

Epoch 11/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0816 - acc: 0.9751 - val_loss: 0.0727 - val_acc: 0.9778

Epoch 12/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0795 - acc: 0.9761 - val_loss: 0.0691 - val_acc: 0.9799

Epoch 13/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0761 - acc: 0.9766 - val_loss: 0.0670 - val_acc: 0.9799

Epoch 14/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0684 - acc: 0.9792 - val_loss: 0.0659 - val_acc: 0.9800

Epoch 15/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0652 - acc: 0.9793 - val_loss: 0.0674 - val_acc: 0.9807

Epoch 16/20

60000/60000 [=====] - 10s 161us/step - loss: 0.0631 - acc: 0.9801 - val_loss: 0.0662 - val_acc: 0.9812

Epoch 17/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0575 - acc: 0.9820 - val_loss: 0.0635 - val_acc: 0.9819

Epoch 18/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0574 - acc: 0.9819 - val_loss: 0.0643 - val_acc: 0.9805

```

val_loss: 0.0643 - val_acc: 0.9803
Epoch 19/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0553 - acc: 0.9826 -
val_loss: 0.0682 - val_acc: 0.9813
Epoch 20/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0513 - acc: 0.9840 -
val_loss: 0.0617 - val_acc: 0.9820

```

In [141]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)


```

Test score: 0.06166264152140356

Test accuracy: 0.982

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).

max_open_warning, RuntimeWarning)



In [142]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```


```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



13. Model 5: MLP + BatchNormalization + Dropout (0.30)

- #layers: 5
- activation: sigmoid
- Weight Initializer: RandomNormal
- Optimizer: adadelata

In [143]:

```

model_relu = Sequential()
model_relu.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.072, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.102, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.144, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adadelata', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_179 (Dense) | (None, 512) | 401920 |
| batch_normalization_74 (Batch Normalization) | (None, 512) | 2048 |
| dropout_57 (Dropout) | (None, 512) | 0 |
| dense_180 (Dense) | (None, 256) | 131328 |
| batch_normalization_75 (Batch Normalization) | (None, 256) | 1024 |
| dropout_58 (Dropout) | (None, 256) | 0 |
| dense_181 (Dense) | (None, 128) | 32896 |
| batch_normalization_76 (Batch Normalization) | (None, 128) | 512 |
| dropout_59 (Dropout) | (None, 128) | 0 |
| dense_182 (Dense) | (None, 64) | 8256 |
| batch_normalization_77 (Batch Normalization) | (None, 64) | 256 |
| dense_183 (Dense) | (None, 32) | 2080 |
| batch_normalization_78 (Batch Normalization) | (None, 32) | 128 |
| dense_184 (Dense) | (None, 10) | 330 |

Total params: 580,778
 Trainable params: 578,794
 Non-trainable params: 1,984

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 19s 316us/step - loss: 0.4114 - acc: 0.8797 - val_loss: 0.2380 - val_acc: 0.9287

Epoch 2/20

60000/60000 [=====] - 9s 156us/step - loss: 0.2509 - acc: 0.9246 - val_loss: 0.1653 - val_acc: 0.9520

Epoch 3/20

60000/60000 [=====] - 9s 157us/step - loss: 0.2036 - acc: 0.9381 - val_loss: 0.1451 - val_acc: 0.9578

Epoch 4/20

60000/60000 [=====] - 9s 154us/step - loss: 0.1751 - acc: 0.9468 - val_loss: 0.1324 - val_acc: 0.9632

Epoch 5/20

60000/60000 [=====] - 9s 156us/step - loss: 0.1528 - acc: 0.9545 - val_loss: 0.1179 - val_acc: 0.9650

Epoch 6/20

60000/60000 [=====] - 9s 156us/step - loss: 0.1415 - acc: 0.9574 - val_loss: 0.1063 - val_acc: 0.9698

Epoch 7/20

60000/60000 [=====] - 9s 154us/step - loss: 0.1299 - acc: 0.9610 - val_loss: 0.0977 - val_acc: 0.9719

Epoch 8/20

60000/60000 [=====] - 9s 155us/step - loss: 0.1196 - acc: 0.9635 - val_loss: 0.0920 - val_acc: 0.9732

Epoch 9/20

60000/60000 [=====] - 9s 155us/step - loss: 0.1147 - acc: 0.9656 - val_loss: 0.0883 - val_acc: 0.9744

Epoch 10/20

60000/60000 [=====] - 9s 155us/step - loss: 0.1083 - acc: 0.9668 - val_loss: 0.0860 - val_acc: 0.9748

Epoch 11/20

60000/60000 [=====] - 9s 155us/step - loss: 0.0996 - acc: 0.9697 - val_loss: 0.0860 - val_acc: 0.9753

Epoch 12/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0948 - acc: 0.9703 - val_loss: 0.0866 - val_acc: 0.9762

Epoch 13/20

60000/60000 [=====] - 9s 155us/step - loss: 0.0896 - acc: 0.9730 - val_loss: 0.0821 - val_acc: 0.9782

```

Epoch 14/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0857 - acc: 0.9733 -
val_loss: 0.0798 - val_acc: 0.9774
Epoch 15/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0828 - acc: 0.9744 -
val_loss: 0.0741 - val_acc: 0.9786
Epoch 16/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0790 - acc: 0.9758 -
val_loss: 0.0696 - val_acc: 0.9793
Epoch 17/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0752 - acc: 0.9759 -
val_loss: 0.0695 - val_acc: 0.9803
Epoch 18/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0721 - acc: 0.9767 -
val_loss: 0.0684 - val_acc: 0.9796
Epoch 19/20
60000/60000 [=====] - 9s 157us/step - loss: 0.0682 - acc: 0.9791 -
val_loss: 0.0729 - val_acc: 0.9788
Epoch 20/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0675 - acc: 0.9789 -
val_loss: 0.0698 - val_acc: 0.9804

```

In [144]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)


```


Test score: 0.06982594733461737

Test accuracy: 0.9804

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).

max_open_warning, RuntimeWarning)





In [145]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



14. Model 6: MLP + BatchNormalization + Dropout (0.30)

- #layers: 5
- activation: tanh
- Weight Initializer: glorot_normal
- Optimizer: ADAM

In [146]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='tanh', input_shape=(input_dim,), kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='tanh', kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='tanh', kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='tanh', kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
```

```

model_relu.add(Dense(32, activation='tanh', kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| ===== | | |
| dense_185 (Dense) | (None, 512) | 401920 |
| batch_normalization_79 (Batch Normalization) | (None, 512) | 2048 |
| dropout_60 (Dropout) | (None, 512) | 0 |
| dense_186 (Dense) | (None, 256) | 131328 |
| batch_normalization_80 (Batch Normalization) | (None, 256) | 1024 |
| dropout_61 (Dropout) | (None, 256) | 0 |
| dense_187 (Dense) | (None, 128) | 32896 |
| batch_normalization_81 (Batch Normalization) | (None, 128) | 512 |
| dropout_62 (Dropout) | (None, 128) | 0 |
| dense_188 (Dense) | (None, 64) | 8256 |
| batch_normalization_82 (Batch Normalization) | (None, 64) | 256 |
| dense_189 (Dense) | (None, 32) | 2080 |
| batch_normalization_83 (Batch Normalization) | (None, 32) | 128 |
| dense_190 (Dense) | (None, 10) | 330 |
| ===== | | |
| Total params: 580,778 | | |
| Trainable params: 578,794 | | |
| Non-trainable params: 1,984 | | |

```

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 19s 312us/step - loss: 0.4344 - acc: 0.8725 - val_loss: 0.2088 - val_acc: 0.9388
Epoch 2/20
60000/60000 [=====] - 9s 145us/step - loss: 0.2380 - acc: 0.9296 - val_loss: 0.1496 - val_acc: 0.9580
Epoch 3/20
60000/60000 [=====] - 9s 146us/step - loss: 0.1820 - acc: 0.9455 - val_loss: 0.1261 - val_acc: 0.9636
Epoch 4/20
60000/60000 [=====] - 9s 158us/step - loss: 0.1496 - acc: 0.9563 - val_loss: 0.1042 - val_acc: 0.9712
Epoch 5/20
60000/60000 [=====] - 9s 149us/step - loss: 0.1287 - acc: 0.9619 - val_loss: 0.1016 - val_acc: 0.9722
Epoch 6/20
60000/60000 [=====] - 9s 146us/step - loss: 0.1195 - acc: 0.9648 - val_loss: 0.0928 - val_acc: 0.9740
Epoch 7/20
60000/60000 [=====] - 9s 146us/step - loss: 0.1040 - acc: 0.9697 - val_loss: 0.0841 - val_acc: 0.9737
Epoch 8/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0968 - acc: 0.9712 - val_loss: 0.0780 - val_acc: 0.9763
Epoch 9/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0900 - acc: 0.9734 - val_loss: 0.0769 - val_acc: 0.9776
Epoch 10/20

```

```

60000/60000 [=====] - 9s 146us/step - loss: 0.0818 - acc: 0.9751 -
val_loss: 0.0716 - val_acc: 0.9796
Epoch 11/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0744 - acc: 0.9770 -
val_loss: 0.0713 - val_acc: 0.9801
Epoch 12/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0755 - acc: 0.9772 -
val_loss: 0.0691 - val_acc: 0.9818
Epoch 13/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0677 - acc: 0.9797 -
val_loss: 0.0741 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0613 - acc: 0.9815 -
val_loss: 0.0655 - val_acc: 0.9832
Epoch 15/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0602 - acc: 0.9821 -
val_loss: 0.0731 - val_acc: 0.9819
Epoch 16/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0596 - acc: 0.9817 -
val_loss: 0.0713 - val_acc: 0.9821
Epoch 17/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0548 - acc: 0.9834 -
val_loss: 0.0730 - val_acc: 0.9814
Epoch 18/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0562 - acc: 0.9825 -
val_loss: 0.0634 - val_acc: 0.9829
Epoch 19/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0508 - acc: 0.9847 -
val_loss: 0.0717 - val_acc: 0.9822
Epoch 20/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0498 - acc: 0.9849 -
val_loss: 0.0628 - val_acc: 0.9843

```

In [147]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```


Test score: 0.06275533262640237
Test accuracy: 0.9843

```

```

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20
figures have been opened. Figures created through the pyplot interface
(`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory.
(To control this warning, see the rcParam `figure.max_open_warning`).
max_open_warning, RuntimeWarning)

```



In [148]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Feedback Conclusion:

- I have used Kaggle platform to do this assignment as I found that kaggle is much much faster than Google colab.
- I have trained MLP Models with 2, 3 and 5 layers.
- I have used RandomNormal, He Normal and Glorot Normal weight initialization.
- I have used ReLU, sigmoid and tanh activation function.
- I have used AdaDelta and ADAM as optimizer.
- ADAM is faster than AdaDelta.
- I have also used BatchNormalization and Dropout.

In [149]:

```
table = PrettyTable()
table.field_names = ['Model #', 'Batch Normalization', 'Dropout + Value', 'Activation', 'Initializer', 'Optimizer', 'Accuracy']
table.add_row([1, "Yes", "Yes, 0.3", "ReLU", "RandomNormal", "ADAM", 0.9832])
table.add_row([2, "No", "Yes, 0.3", "ReLU", "He Normal", "ADAM", 0.9833])
table.add_row([3, "Yes", "Yes, 0.4", "ReLU", "RandomNormal", "ADAM", 0.9836])
table.add_row([4, "Yes", "Yes, 0.3", "sigmoid", "RandomNormal", "ADAM", 0.9820])
table.add_row([5, "Yes", "Yes, 0.3", "sigmoid", "RandomNormal", "AdaDelta", 0.9804])
table.add_row([6, "Yes", "Yes, 0.3", "tanh", "Glorot Normal", "ADAM", 0.9848])
print(table)
```

| Model # | Batch Normalization | Dropout + Value | Activation | Initializer | Optimizer | Accuracy |
|---------|---------------------|-----------------|------------|---------------|-----------|----------|
| 1 | Yes | Yes, 0.3 | ReLU | RandomNormal | ADAM | 0.9832 |
| 2 | No | Yes, 0.3 | ReLU | He Normal | ADAM | 0.9833 |
| 3 | Yes | Yes, 0.4 | ReLU | RandomNormal | ADAM | 0.9836 |
| 4 | Yes | Yes, 0.3 | sigmoid | RandomNormal | ADAM | 0.9820 |
| 5 | Yes | Yes, 0.3 | sigmoid | RandomNormal | AdaDelta | 0.9804 |
| 6 | Yes | Yes, 0.3 | tanh | Glorot Normal | ADAM | 0.9848 |