

Keras -- MLPs on MNIST

In [80]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

In [81]:

```
from keras.initializers import glorot_normal
from keras.initializers import he_normal
```

In [82]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [83]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [84]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of test examples : 10000 and each image is of shape (28, 28)

In [85]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [86]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of test examples : 10000 and each image is of shape (784)

In [87]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

In [88]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
X_train = X_train/255
X_test = X_test/255
```

In [89]:

```
# example data point after normlizing
print(X_train[0])
```

```
[0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.]
```

[illegible]

[illegible]

In [90]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

In [91]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor.
```

```

# you can create a sequential model by passing a list of layer instances to the constructor.

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [92]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

In [93]:

```

# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

```

```
model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

In [94]:

```
# Before training a model, you need to configure the learning process, which is done via the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer ,
https://keras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize.,
https://keras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=['accuracy'].
https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categorical format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that
is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 111us/step - loss: 1.2659 - acc: 0.7094 -
val_loss: 0.8067 - val_acc: 0.8280
Epoch 2/20
60000/60000 [=====] - 3s 42us/step - loss: 0.7120 - acc: 0.8417 -
val_loss: 0.6061 - val_acc: 0.8616
Epoch 3/20
60000/60000 [=====] - 3s 43us/step - loss: 0.5854 - acc: 0.8597 -
val_loss: 0.5252 - val_acc: 0.8740
Epoch 4/20
60000/60000 [=====] - 3s 43us/step - loss: 0.5244 - acc: 0.8694 -
val_loss: 0.4799 - val_acc: 0.8815
Epoch 5/20
60000/60000 [=====] - 3s 42us/step - loss: 0.4871 - acc: 0.8753 -
val_loss: 0.4507 - val_acc: 0.8866
Epoch 6/20
60000/60000 [=====] - 3s 43us/step - loss: 0.4614 - acc: 0.8796 -
val_loss: 0.4292 - val_acc: 0.8889
Epoch 7/20
60000/60000 [=====] - 3s 42us/step - loss: 0.4423 - acc: 0.8833 -
val_loss: 0.4131 - val_acc: 0.8920
Epoch 8/20
60000/60000 [=====] - 3s 44us/step - loss: 0.4274 - acc: 0.8864 -
val_loss: 0.4002 - val_acc: 0.8949
Epoch 9/20
60000/60000 [=====] - 3s 43us/step - loss: 0.4154 - acc: 0.8886 -
val loss: 0.3901 - val acc: 0.8962
```

```

Epoch 10/20
60000/60000 [=====] - 3s 43us/step - loss: 0.4054 - acc: 0.8904 -
val_loss: 0.3811 - val_acc: 0.8976
Epoch 11/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3969 - acc: 0.8922 -
val_loss: 0.3739 - val_acc: 0.9004
Epoch 12/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3896 - acc: 0.8939 -
val_loss: 0.3671 - val_acc: 0.9017
Epoch 13/20
60000/60000 [=====] - 3s 43us/step - loss: 0.3833 - acc: 0.8952 -
val_loss: 0.3617 - val_acc: 0.9027
Epoch 14/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3777 - acc: 0.8965 -
val_loss: 0.3568 - val_acc: 0.9033
Epoch 15/20
60000/60000 [=====] - 3s 43us/step - loss: 0.3726 - acc: 0.8975 -
val_loss: 0.3524 - val_acc: 0.9046
Epoch 16/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3681 - acc: 0.8986 -
val_loss: 0.3481 - val_acc: 0.9047
Epoch 17/20
60000/60000 [=====] - 2s 41us/step - loss: 0.3640 - acc: 0.8994 -
val_loss: 0.3447 - val_acc: 0.9052
Epoch 18/20
60000/60000 [=====] - 3s 43us/step - loss: 0.3603 - acc: 0.9003 -
val_loss: 0.3414 - val_acc: 0.9069
Epoch 19/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3568 - acc: 0.9013 -
val_loss: 0.3383 - val_acc: 0.9076
Epoch 20/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3537 - acc: 0.9020 -
val_loss: 0.3357 - val_acc: 0.9086

```

In [95]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.3357250750780106
Test accuracy: 0.9086





MLP + Sigmoid activation + SGDOptimizer

In [96]:

```
# Multilayer perceptron

model_sigmoind = Sequential()
model_sigmoind.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoind.add(Dense(128, activation='sigmoid'))
model_sigmoind.add(Dense(output_dim, activation='softmax'))

model_sigmoind.summary()
```

Layer (type)	Output Shape	Param #
dense_103 (Dense)	(None, 512)	401920
dense_104 (Dense)	(None, 128)	65664
dense_105 (Dense)	(None, 10)	1290

Total params: 468,874
 Trainable params: 468,874
 Non-trainable params: 0

In [97]:

```
model_sigmoind.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoind.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 116us/step - loss: 2.2665 - acc: 0.2311 -
val_loss: 2.2210 - val_acc: 0.4829
Epoch 2/20
60000/60000 [=====] - 3s 46us/step - loss: 2.1767 - acc: 0.4668 -
val_loss: 2.1211 - val_acc: 0.5636
Epoch 3/20
60000/60000 [=====] - 3s 47us/step - loss: 2.0609 - acc: 0.5972 -
val_loss: 1.9797 - val_acc: 0.6697
Epoch 4/20
60000/60000 [=====] - 3s 46us/step - loss: 1.8947 - acc: 0.6608 -
val_loss: 1.7812 - val_acc: 0.6553
Epoch 5/20
60000/60000 [=====] - 3s 47us/step - loss: 1.6747 - acc: 0.6968 -
val_loss: 1.5382 - val_acc: 0.7561
Epoch 6/20
60000/60000 [=====] - 3s 47us/step - loss: 1.4344 - acc: 0.7361 -
val_loss: 1.3019 - val_acc: 0.7611
Epoch 7/20
60000/60000 [=====] - 3s 46us/step - loss: 1.2204 - acc: 0.7636 -
val_loss: 1.1097 - val_acc: 0.7879
Epoch 8/20
60000/60000 [=====] - 3s 46us/step - loss: 1.0534 - acc: 0.7845 -
val_loss: 0.9659 - val_acc: 0.7938
Epoch 9/20
60000/60000 [=====] - 3s 46us/step - loss: 0.9287 - acc: 0.7973 -
val_loss: 0.8579 - val_acc: 0.8082
```



```

Epoch 10/20
60000/60000 [=====] - 3s 45us/step - loss: 0.8347 - acc: 0.8097 -
val_loss: 0.7763 - val_acc: 0.8234
Epoch 11/20
60000/60000 [=====] - 3s 46us/step - loss: 0.7623 - acc: 0.8208 -
val_loss: 0.7125 - val_acc: 0.8313
Epoch 12/20
60000/60000 [=====] - 3s 46us/step - loss: 0.7049 - acc: 0.8292 -
val_loss: 0.6622 - val_acc: 0.8421
Epoch 13/20
60000/60000 [=====] - 3s 45us/step - loss: 0.6585 - acc: 0.8369 -
val_loss: 0.6196 - val_acc: 0.8463
Epoch 14/20
60000/60000 [=====] - 3s 46us/step - loss: 0.6206 - acc: 0.8437 -
val_loss: 0.5854 - val_acc: 0.8509
Epoch 15/20
60000/60000 [=====] - 3s 46us/step - loss: 0.5888 - acc: 0.8493 -
val_loss: 0.5572 - val_acc: 0.8578
Epoch 16/20
60000/60000 [=====] - 3s 45us/step - loss: 0.5619 - acc: 0.8549 -
val_loss: 0.5329 - val_acc: 0.8620
Epoch 17/20
60000/60000 [=====] - 3s 45us/step - loss: 0.5388 - acc: 0.8593 -
val_loss: 0.5110 - val_acc: 0.8666
Epoch 18/20
60000/60000 [=====] - 3s 46us/step - loss: 0.5187 - acc: 0.8634 -
val_loss: 0.4927 - val_acc: 0.8702
Epoch 19/20
60000/60000 [=====] - 3s 46us/step - loss: 0.5012 - acc: 0.8671 -
val_loss: 0.4761 - val_acc: 0.8727
Epoch 20/20
60000/60000 [=====] - 3s 46us/step - loss: 0.4856 - acc: 0.8700 -
val_loss: 0.4625 - val_acc: 0.8764

```

In [98]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.4624546233654022

Test accuracy: 0.8764





In [99]:


```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Sigmoid activation + ADAM

In [100]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_106 (Dense)	(None, 512)	401920
dense_107 (Dense)	(None, 128)	65664
dense_108 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 125us/step - loss: 0.5394 - acc: 0.8597 -
val_loss: 0.2513 - val_acc: 0.9265
Epoch 2/20
60000/60000 [=====] - 3s 53us/step - loss: 0.2212 - acc: 0.9357 -
val_loss: 0.1897 - val_acc: 0.9416
Epoch 3/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1631 - acc: 0.9518 -
val_loss: 0.1475 - val_acc: 0.9547
Epoch 4/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1251 - acc: 0.9627 -
val_loss: 0.1142 - val_acc: 0.9649
Epoch 5/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0985 - acc: 0.9711 -
val_loss: 0.0977 - val_acc: 0.9703
Epoch 6/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0793 - acc: 0.9762 -
val_loss: 0.0865 - val_acc: 0.9729
Epoch 7/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0631 - acc: 0.9812 -
val_loss: 0.0771 - val_acc: 0.9760
Epoch 8/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0514 - acc: 0.9846 -
val_loss: 0.0805 - val_acc: 0.9754
Epoch 9/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0421 - acc: 0.9879 -
val_loss: 0.0718 - val_acc: 0.9781
Epoch 10/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0337 - acc: 0.9906 -
val_loss: 0.0650 - val_acc: 0.9789
Epoch 11/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0268 - acc: 0.9929 -
val_loss: 0.0655 - val_acc: 0.9803
Epoch 12/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0223 - acc: 0.9940 -
val_loss: 0.0668 - val_acc: 0.9797
Epoch 13/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0175 - acc: 0.9956 -
val_loss: 0.0628 - val_acc: 0.9813
Epoch 14/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0141 - acc: 0.9966 -
val_loss: 0.0650 - val_acc: 0.9818
Epoch 15/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0119 - acc: 0.9971 -
val_loss: 0.0682 - val_acc: 0.9802
Epoch 16/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0084 - acc: 0.9983 -
val_loss: 0.0679 - val_acc: 0.9812
Epoch 17/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0075 - acc: 0.9984 -
val_loss: 0.0651 - val_acc: 0.9819
Epoch 18/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0056 - acc: 0.9990 -
val_loss: 0.0647 - val_acc: 0.9823
Epoch 19/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0053 - acc: 0.9988 -
val_loss: 0.0727 - val_acc: 0.9809
Epoch 20/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0034 - acc: 0.9994 -
val_loss: 0.0700 - val_acc: 0.9826
```

In [101]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lilation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0699703871981983
Test accuracy: 0.9826



In [102]:

```

w_after = model_sigmoid.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + ReLU +SGD

In [103]:

```
# Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_109 (Dense)	(None, 512)	401920
dense_110 (Dense)	(None, 128)	65664
dense_111 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [104]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 120us/step - loss: 0.7727 - acc: 0.7796 - val_loss: 0.3984 - val_acc: 0.8890
Epoch 2/20
60000/60000 [=====] - 3s 47us/step - loss: 0.3619 - acc: 0.8984 - val_loss: 0.3029 - val_acc: 0.9131
Epoch 3/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3245 - acc: 0.9160 - val_loss: 0.2845 - val_acc: 0.9210
```

```

60000/60000 [=====] - 3s 48us/step - loss: 0.2946 - acc: 0.9168 -
val_loss: 0.2617 - val_acc: 0.9263
Epoch 4/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2585 - acc: 0.9263 -
val_loss: 0.2349 - val_acc: 0.9343
Epoch 5/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2330 - acc: 0.9342 -
val_loss: 0.2166 - val_acc: 0.9387
Epoch 6/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2137 - acc: 0.9399 -
val_loss: 0.2030 - val_acc: 0.9409
Epoch 7/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1983 - acc: 0.9437 -
val_loss: 0.1905 - val_acc: 0.9453
Epoch 8/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1855 - acc: 0.9478 -
val_loss: 0.1810 - val_acc: 0.9477
Epoch 9/20
60000/60000 [=====] - 3s 46us/step - loss: 0.1747 - acc: 0.9501 -
val_loss: 0.1721 - val_acc: 0.9485
Epoch 10/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1651 - acc: 0.9534 -
val_loss: 0.1646 - val_acc: 0.9510
Epoch 11/20
60000/60000 [=====] - 3s 45us/step - loss: 0.1567 - acc: 0.9556 -
val_loss: 0.1579 - val_acc: 0.9527
Epoch 12/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1494 - acc: 0.9573 -
val_loss: 0.1536 - val_acc: 0.9545
Epoch 13/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1424 - acc: 0.9602 -
val_loss: 0.1474 - val_acc: 0.9564
Epoch 14/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1366 - acc: 0.9618 -
val_loss: 0.1424 - val_acc: 0.9571
Epoch 15/20
60000/60000 [=====] - 3s 49us/step - loss: 0.1311 - acc: 0.9636 -
val_loss: 0.1413 - val_acc: 0.9574
Epoch 16/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1259 - acc: 0.9649 -
val_loss: 0.1350 - val_acc: 0.9590
Epoch 17/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1211 - acc: 0.9660 -
val_loss: 0.1317 - val_acc: 0.9595
Epoch 18/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1166 - acc: 0.9677 -
val_loss: 0.1276 - val_acc: 0.9609
Epoch 19/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1125 - acc: 0.9689 -
val_loss: 0.1260 - val_acc: 0.9610
Epoch 20/20
60000/60000 [=====] - 3s 46us/step - loss: 0.1087 - acc: 0.9698 -
val_loss: 0.1225 - val_acc: 0.9627

```

In [105]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss

```

```
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12253610298782587
Test accuracy: 0.9627



In [106]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU + ADAM

In [107]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_112 (Dense)	(None, 512)	401920
dense_113 (Dense)	(None, 128)	65664
dense_114 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 130us/step - loss: 0.2288 - acc: 0.9326 - val_loss: 0.1132 - val_acc: 0.9655
Epoch 2/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0854 - acc: 0.9737 - val_loss: 0.0770 - val_acc: 0.9761
Epoch 3/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0522 - acc: 0.9837 - val_loss: 0.0829 - val_acc: 0.9741
Epoch 4/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0380 - acc: 0.9879 - val_loss: 0.0755 - val_acc: 0.9764
Epoch 5/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0262 - acc: 0.9918 - val_loss: 0.0712 - val_acc: 0.9793
Epoch 6/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0200 - acc: 0.9940 - val_loss: 0.0768 - val_acc: 0.9788
Epoch 7/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0152 - acc: 0.9951 - val_loss: 0.0642 - val_acc: 0.9828
Epoch 8/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0151 - acc: 0.9950 - val_loss: 0.0781 - val_acc: 0.9796
Epoch 9/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0151 - acc: 0.9950 - val_loss: 0.0795 - val_acc: 0.9795
Epoch 10/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0111 - acc: 0.9962 - val_loss: 0.0853 - val_acc: 0.9806
Epoch 11/20
60000/60000 [=====] - 3s 54us/step - loss: 0.0105 - acc: 0.9965 - val_loss: 0.1058 - val_acc: 0.9755
Epoch 12/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0129 - acc: 0.9956 - val_loss: 0.0834 - val_acc: 0.9799
Epoch 13/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0072 - acc: 0.9977 - val_loss: 0.0840 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0095 - acc: 0.9968 -


```

val_loss: 0.0819 - val_acc: 0.9814
Epoch 15/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0096 - acc: 0.9967 -
val_loss: 0.1173 - val_acc: 0.9756
Epoch 16/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0096 - acc: 0.9966 -
val_loss: 0.0950 - val_acc: 0.9800
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0068 - acc: 0.9978 -
val_loss: 0.1090 - val_acc: 0.9777
Epoch 18/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0093 - acc: 0.9972 -
val_loss: 0.1086 - val_acc: 0.9780
Epoch 19/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0082 - acc: 0.9971 -
val_loss: 0.0813 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0038 - acc: 0.9988 -
val_loss: 0.0990 - val_acc: 0.9826

```

In [108]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09899708161438571

Test accuracy: 0.9826



In [109]:

```

w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

In [110]:

```

# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_115 (Dense)	(None, 512)	401920
batch_normalization_45 (Batch Normalization)	(None, 512)	2048
dense_116 (Dense)	(None, 128)	65664
batch_normalization_46 (Batch Normalization)	(None, 128)	512
dense_117 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [111]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 10s 163us/step - loss: 0.3013 - acc: 0.9113 - val_loss: 0.2207 - val_acc: 0.9347
Epoch 2/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1757 - acc: 0.9487 - val_loss: 0.1698 - val_acc: 0.9502
Epoch 3/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1382 - acc: 0.9593 - val_loss: 0.1542 - val_acc: 0.9546
Epoch 4/20
60000/60000 [=====] - 5s 80us/step - loss: 0.1140 - acc: 0.9656 - val_loss: 0.1405 - val_acc: 0.9581
Epoch 5/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0947 - acc: 0.9720 - val_loss: 0.1276 - val_acc: 0.9625
Epoch 6/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0804 - acc: 0.9753 - val_loss: 0.1157 - val_acc: 0.9654
Epoch 7/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0674 - acc: 0.9795 - val_loss: 0.1131 - val_acc: 0.9656
Epoch 8/20
60000/60000 [=====] - 5s 84us/step - loss: 0.0582 - acc: 0.9826 - val_loss: 0.1060 - val_acc: 0.9675
Epoch 9/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0509 - acc: 0.9842 - val_loss: 0.1078 - val_acc: 0.9687
Epoch 10/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0474 - acc: 0.9853 - val_loss: 0.1032 - val_acc: 0.9681
Epoch 11/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0395 - acc: 0.9876 - val_loss: 0.0999 - val_acc: 0.9705
Epoch 12/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0341 - acc: 0.9889 - val_loss: 0.0963 - val_acc: 0.9719
Epoch 13/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0308 - acc: 0.9902 - val_loss: 0.0936 - val_acc: 0.9710
Epoch 14/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0265 - acc: 0.9917 - val_loss: 0.0916 - val_acc: 0.9733
Epoch 15/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0224 - acc: 0.9933 - val_loss: 0.1015 - val_acc: 0.9717
Epoch 16/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0202 - acc: 0.9936 - val_loss: 0.0969 - val_acc: 0.9737
Epoch 17/20
```

```
Epoch 17/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0208 - acc: 0.9932 -
val_loss: 0.1138 - val_acc: 0.9698
Epoch 18/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0196 - acc: 0.9936 -
val_loss: 0.0960 - val_acc: 0.9750
Epoch 19/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0177 - acc: 0.9941 -
val_loss: 0.0924 - val_acc: 0.9749
Epoch 20/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0169 - acc: 0.9944 -
val_loss: 0.0983 - val_acc: 0.9734
```

In [112]:

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09825632937119808

Test accuracy: 0.9734



In [113]:

```
w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
```

```
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5. MLP + Dropout + AdamOptimizer

In [114]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_118 (Dense)	(None, 512)	401920
batch_normalization_47 (Batch Normalization)	(None, 512)	2048
dropout_35 (Dropout)	(None, 512)	0
dense_119 (Dense)	(None, 128)	65664
batch_normalization_48 (Batch Normalization)	(None, 128)	512

dropout_36 (Dropout)	(None, 128)	0
dense_120 (Dense)	(None, 10)	1290
=====		
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		
=====		

In [115]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

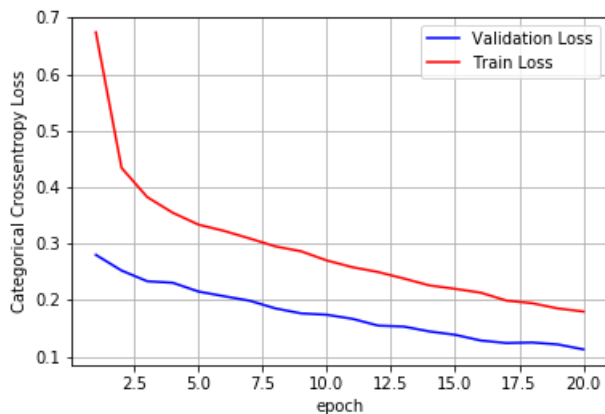
```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 177us/step - loss: 0.6745 - acc: 0.7926 - val_loss: 0.2803 - val_acc: 0.9194
Epoch 2/20
60000/60000 [=====] - 5s 85us/step - loss: 0.4344 - acc: 0.8685 - val_loss: 0.2526 - val_acc: 0.9237
Epoch 3/20
60000/60000 [=====] - 5s 84us/step - loss: 0.3827 - acc: 0.8839 - val_loss: 0.2336 - val_acc: 0.9312
Epoch 4/20
60000/60000 [=====] - 6s 98us/step - loss: 0.3548 - acc: 0.8927 - val_loss: 0.2309 - val_acc: 0.9323
Epoch 5/20
60000/60000 [=====] - 5s 86us/step - loss: 0.3339 - acc: 0.8993 - val_loss: 0.2153 - val_acc: 0.9359
Epoch 6/20
60000/60000 [=====] - 5s 84us/step - loss: 0.3227 - acc: 0.9015 - val_loss: 0.2070 - val_acc: 0.9396
Epoch 7/20
60000/60000 [=====] - 5s 85us/step - loss: 0.3092 - acc: 0.9063 - val_loss: 0.1990 - val_acc: 0.9406
Epoch 8/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2952 - acc: 0.9118 - val_loss: 0.1854 - val_acc: 0.9429
Epoch 9/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2866 - acc: 0.9136 - val_loss: 0.1767 - val_acc: 0.9462
Epoch 10/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2704 - acc: 0.9189 - val_loss: 0.1742 - val_acc: 0.9467
Epoch 11/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2585 - acc: 0.9223 - val_loss: 0.1670 - val_acc: 0.9505
Epoch 12/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2499 - acc: 0.9249 - val_loss: 0.1550 - val_acc: 0.9541
Epoch 13/20
60000/60000 [=====] - 5s 84us/step - loss: 0.2384 - acc: 0.9281 - val_loss: 0.1533 - val_acc: 0.9534
Epoch 14/20
60000/60000 [=====] - 5s 84us/step - loss: 0.2262 - acc: 0.9318 - val_loss: 0.1447 - val_acc: 0.9560
Epoch 15/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2201 - acc: 0.9342 - val_loss: 0.1387 - val_acc: 0.9570
Epoch 16/20
60000/60000 [=====] - 5s 85us/step - loss: 0.2133 - acc: 0.9354 - val_loss: 0.1286 - val_acc: 0.9619
Epoch 17/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1993 - acc: 0.9406 - val_loss: 0.1242 - val_acc: 0.9622
Epoch 18/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1947 - acc: 0.9415 - val_loss: 0.1251 - val_acc: 0.9626
Epoch 19/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1854 - acc: 0.9446 - val_loss: 0.1217 - val_acc: 0.9649
Epoch 20/20
```

```
Epoch 20/20  
60000/60000 [=====] - 5s 84us/step - loss: 0.1799 - acc: 0.9449 -  
val_loss: 0.1129 - val_acc: 0.9685
```

In [116]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])  
  
fig, ax = plt.subplots(1, 1)  
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')  
  
# list of epoch numbers  
x = list(range(1, nb_epoch+1))  
  
# print(history.history.keys())  
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])  
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va  
# lidaion_data=(X_test, Y_test))  
  
# we will get val_loss and val_acc only when you pass the paramter validation_data  
# val_loss : validation loss  
# val_acc : validation accuracy  
  
# loss : training loss  
# acc : train accuracy  
# for each key in history.history we will have a list of length equal to number of epochs  
  
vy = history.history['val_loss']  
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```


Test score: 0.11289959549605846
Test accuracy: 0.9685



In [117]:

```
w_after = model_drop.get_weights()  
  
h1_w = w_after[0].flatten().reshape(-1,1)  
h2_w = w_after[2].flatten().reshape(-1,1)  
out_w = w_after[4].flatten().reshape(-1,1)  
  
fig = plt.figure()  
plt.title("Weight matrices after model trained")  
plt.subplot(1, 3, 1)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h1_w, color='b')  
plt.xlabel('Hidden Layer 1')  
  
plt.subplot(1, 3, 2)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h2_w, color='r')  
plt.xlabel('Hidden Layer 2')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hyper-parameter tuning of Keras models using Sklearn

In [118]:

```
from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

In [119]:

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid','relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
/opt/conda/lib/python3.6/site-packages/sklearn/model_selection/_split.py:1978: FutureWarning: The default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this warning.
  warnings.warn(CV_WARNING, FutureWarning)
```


In [120]:

```
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.976300 using {'activ': 'relu'}
0.975617 (0.001652) with: {'activ': 'sigmoid'}
0.976300 (0.002729) with: {'activ': 'relu'}
```

Assignment

Model 1 with two hidden layers + Batch Normalization + Dropout:

In [121]:

```
model_relu = Sequential()
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_142 (Dense)	(None, 256)	200960
batch_normalization_49 (Batch Normalization)	(None, 256)	1024
dropout_37 (Dropout)	(None, 256)	0
dense_143 (Dense)	(None, 128)	32896
batch_normalization_50 (Batch Normalization)	(None, 128)	512
dropout_38 (Dropout)	(None, 128)	0
dense_144 (Dense)	(None, 10)	1290

Total params: 236,682
Trainable params: 235,914
Non-trainable params: 768

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 192us/step - loss: 0.3786 - acc: 0.8847 - val_loss: 0.1445 - val_acc: 0.9560
Epoch 2/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1768 - acc: 0.9470 - val_loss: 0.1059 - val_acc: 0.9682
Epoch 3/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1326 - acc: 0.9583 - val_loss: 0.0858 - val_acc: 0.9728
Epoch 4/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1104 - acc: 0.9652 -

```

val_loss: 0.0835 - val_acc: 0.9741
Epoch 5/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0954 - acc: 0.9700 -
val_loss: 0.0715 - val_acc: 0.9782
Epoch 6/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0858 - acc: 0.9732 -
val_loss: 0.0671 - val_acc: 0.9801
Epoch 7/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0755 - acc: 0.9759 -
val_loss: 0.0704 - val_acc: 0.9786
Epoch 8/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0683 - acc: 0.9783 -
val_loss: 0.0650 - val_acc: 0.9805
Epoch 9/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0652 - acc: 0.9794 -
val_loss: 0.0627 - val_acc: 0.9810
Epoch 10/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0610 - acc: 0.9803 -
val_loss: 0.0710 - val_acc: 0.9803
Epoch 11/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0559 - acc: 0.9811 -
val_loss: 0.0630 - val_acc: 0.9816
Epoch 12/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0525 - acc: 0.9830 -
val_loss: 0.0640 - val_acc: 0.9814
Epoch 13/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0500 - acc: 0.9840 -
val_loss: 0.0710 - val_acc: 0.9798
Epoch 14/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0459 - acc: 0.9850 -
val_loss: 0.0672 - val_acc: 0.9807
Epoch 15/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0456 - acc: 0.9844 -
val_loss: 0.0621 - val_acc: 0.9820
Epoch 16/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0428 - acc: 0.9859 -
val_loss: 0.0575 - val_acc: 0.9826
Epoch 17/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0427 - acc: 0.9861 -
val_loss: 0.0585 - val_acc: 0.9826
Epoch 18/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0401 - acc: 0.9862 -
val_loss: 0.0611 - val_acc: 0.9833
Epoch 19/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0387 - acc: 0.9876 -
val_loss: 0.0600 - val_acc: 0.9818
Epoch 20/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0358 - acc: 0.9880 -
val_loss: 0.0637 - val_acc: 0.9827

```

In [122]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06367886031113303
Test accuracy: 0.9827



In [123]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Model 2 with three hidden lavers + Batch Normalization + Dropout:

In [124]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====		
dense_145 (Dense)	(None, 512)	401920
batch_normalization_51 (Batch Normalization)	(None, 512)	2048
dropout_39 (Dropout)	(None, 512)	0
dense_146 (Dense)	(None, 256)	131328
batch_normalization_52 (Batch Normalization)	(None, 256)	1024
dropout_40 (Dropout)	(None, 256)	0
dense_147 (Dense)	(None, 128)	32896
batch_normalization_53 (Batch Normalization)	(None, 128)	512
dropout_41 (Dropout)	(None, 128)	0
dense_148 (Dense)	(None, 10)	1290
=====		
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

```
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 13s 224us/step - loss: 0.3627 - acc: 0.8891 - val_loss: 0.1164 - val_acc: 0.9641
Epoch 2/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1567 - acc: 0.9530 - val_loss: 0.0919 - val_acc: 0.9695
Epoch 3/20
60000/60000 [=====] - 6s 108us/step - loss: 0.1184 - acc: 0.9636 - val_loss: 0.0804 - val_acc: 0.9751
Epoch 4/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1020 - acc: 0.9687 - val_loss: 0.0829 - val_acc: 0.9746
Epoch 5/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0844 - acc: 0.9733 - val_loss: 0.0787 - val_acc: 0.9756
Epoch 6/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0770 - acc: 0.9754 - val_loss: 0.0699 - val_acc: 0.9785
Epoch 7/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0689 - acc: 0.9780 - val_loss: 0.0686 - val_acc: 0.9770
Epoch 8/20
```

```

60000/60000 [=====] - 7s 118us/step - loss: 0.0626 - acc: 0.9799 -
val_loss: 0.0690 - val_acc: 0.9794
Epoch 9/20
60000/60000 [=====] - 7s 111us/step - loss: 0.0568 - acc: 0.9819 -
val_loss: 0.0632 - val_acc: 0.9808
Epoch 10/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0511 - acc: 0.9834 -
val_loss: 0.0666 - val_acc: 0.9802
Epoch 11/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0495 - acc: 0.9839 -
val_loss: 0.0640 - val_acc: 0.9800
Epoch 12/20
60000/60000 [=====] - 7s 109us/step - loss: 0.0475 - acc: 0.9845 -
val_loss: 0.0567 - val_acc: 0.9833
Epoch 13/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0444 - acc: 0.9856 -
val_loss: 0.0566 - val_acc: 0.9836
Epoch 14/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0388 - acc: 0.9875 -
val_loss: 0.0595 - val_acc: 0.9823
Epoch 15/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0391 - acc: 0.9874 -
val_loss: 0.0619 - val_acc: 0.9820
Epoch 16/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0378 - acc: 0.9879 -
val_loss: 0.0609 - val_acc: 0.9828
Epoch 17/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0332 - acc: 0.9890 -
val_loss: 0.0576 - val_acc: 0.9844
Epoch 18/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0349 - acc: 0.9884 -
val_loss: 0.0621 - val_acc: 0.9821
Epoch 19/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0323 - acc: 0.9894 -
val_loss: 0.0549 - val_acc: 0.9857
Epoch 20/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0317 - acc: 0.9896 -
val_loss: 0.0609 - val_acc: 0.9843

```

In [125]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lida
tion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.060851101654174275
Test accuracy: 0.9843

```



In [126]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Model 3 with five hidden layers + Batch Normalization + Dropout:

In [127]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
```

```

model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
=====		
dense_149 (Dense)	(None, 512)	401920
batch_normalization_54 (Batch Normalization)	(None, 512)	2048
dropout_42 (Dropout)	(None, 512)	0
dense_150 (Dense)	(None, 256)	131328
batch_normalization_55 (Batch Normalization)	(None, 256)	1024
dropout_43 (Dropout)	(None, 256)	0
dense_151 (Dense)	(None, 128)	32896
batch_normalization_56 (Batch Normalization)	(None, 128)	512
dropout_44 (Dropout)	(None, 128)	0
dense_152 (Dense)	(None, 64)	8256
batch_normalization_57 (Batch Normalization)	(None, 64)	256
dense_153 (Dense)	(None, 32)	2080
batch_normalization_58 (Batch Normalization)	(None, 32)	128
dense_154 (Dense)	(None, 10)	330
=====		

Total params: 580,778

Trainable params: 578,794

Non-trainable params: 1,984

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 16s 273us/step - loss: 0.4450 - acc: 0.8676 - val_loss: 0.1393 - val_acc: 0.9595

Epoch 2/20

60000/60000 [=====] - 9s 144us/step - loss: 0.1794 - acc: 0.9467 - val_loss: 0.1054 - val_acc: 0.9682

Epoch 3/20

60000/60000 [=====] - 9s 143us/step - loss: 0.1373 - acc: 0.9597 - val_loss: 0.0938 - val_acc: 0.9715

Epoch 4/20

60000/60000 [=====] - 9s 144us/step - loss: 0.1131 - acc: 0.9658 - val_loss: 0.0848 - val_acc: 0.9747

Epoch 5/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0990 - acc: 0.9699 - val_loss: 0.0750 - val_acc: 0.9779

Epoch 6/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0866 - acc: 0.9733 - val_loss: 0.0688 - val_acc: 0.9805

Epoch 7/20

```

60000/60000 [=====] - 8s 141us/step - loss: 0.0795 - acc: 0.9754 -
val_loss: 0.0703 - val_acc: 0.9788
Epoch 8/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0718 - acc: 0.9781 -
val_loss: 0.0624 - val_acc: 0.9816
Epoch 9/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0664 - acc: 0.9800 -
val_loss: 0.0672 - val_acc: 0.9805
Epoch 10/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0625 - acc: 0.9800 -
val_loss: 0.0611 - val_acc: 0.9813
Epoch 11/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0588 - acc: 0.9817 -
val_loss: 0.0627 - val_acc: 0.9824
Epoch 12/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0559 - acc: 0.9827 -
val_loss: 0.0598 - val_acc: 0.9830
Epoch 13/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0528 - acc: 0.9834 -
val_loss: 0.0611 - val_acc: 0.9839
Epoch 14/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0510 - acc: 0.9841 -
val_loss: 0.0659 - val_acc: 0.9822
Epoch 15/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0460 - acc: 0.9860 -
val_loss: 0.0639 - val_acc: 0.9822
Epoch 16/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0444 - acc: 0.9861 -
val_loss: 0.0593 - val_acc: 0.9834
Epoch 17/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0433 - acc: 0.9865 -
val_loss: 0.0599 - val_acc: 0.9835
Epoch 18/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0391 - acc: 0.9880 -
val_loss: 0.0582 - val_acc: 0.9850
Epoch 19/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0407 - acc: 0.9875 -
val_loss: 0.0593 - val_acc: 0.9835
Epoch 20/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0376 - acc: 0.9880 -
val_loss: 0.0561 - val_acc: 0.9843

```

In [128]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.05613952756321523
Test accuracy: 0.9843

```

In [129]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Conclusion:

In [150]:

```
from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["# Layers", "Epoch", "Accuracy"]
x.add_row(["2", 20, 0.9827])
```

```
x.add_row(["3", 20, 0.9843])
x.add_row(["5", 20, 0.9843])
print(x)
```

```
+-----+-----+-----+
| # Layers | Epoch | Accuracy |
+-----+-----+-----+
| 2       | 20    | 0.9827   |
| 3       | 20    | 0.9843   |
| 5       | 20    | 0.9843   |
+-----+-----+-----+
```

Assignment Feedback Experimental Models

Model 1: MLP + BatchNormalization + Dropout (0.30)

- #layers: 5
- activation: ReLU
- Weight Initializer: RandomNormal
- Optimizer: ADAM

In [131]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.050, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_155 (Dense)	(None, 512)	401920
batch_normalization_59 (Batch Normalization)	(None, 512)	2048
dropout_45 (Dropout)	(None, 512)	0
dense_156 (Dense)	(None, 256)	131328
batch_normalization_60 (Batch Normalization)	(None, 256)	1024
dropout_46 (Dropout)	(None, 256)	0
dense_157 (Dense)	(None, 128)	32896
batch_normalization_61 (Batch Normalization)	(None, 128)	512

dropout_47 (Dropout)	(None, 128)	0
dense_158 (Dense)	(None, 64)	8256
batch_normalization_62 (Batch Normalization)	(None, 64)	256
dense_159 (Dense)	(None, 32)	2080
batch_normalization_63 (Batch Normalization)	(None, 32)	128
dense_160 (Dense)	(None, 10)	330

=====
Total params: 580,778
Trainable params: 578,794
Non-trainable params: 1,984

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 278us/step - loss: 0.4477 - acc: 0.8703 - val_loss: 0.1462 - val_acc: 0.9558

Epoch 2/20

60000/60000 [=====] - 8s 142us/step - loss: 0.1781 - acc: 0.9478 - val_loss: 0.0991 - val_acc: 0.9717

Epoch 3/20

60000/60000 [=====] - 9s 158us/step - loss: 0.1386 - acc: 0.9588 - val_loss: 0.0854 - val_acc: 0.9741

Epoch 4/20

60000/60000 [=====] - 9s 144us/step - loss: 0.1147 - acc: 0.9652 - val_loss: 0.0748 - val_acc: 0.9775

Epoch 5/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0966 - acc: 0.9700 - val_loss: 0.0822 - val_acc: 0.9767

Epoch 6/20

60000/60000 [=====] - 8s 141us/step - loss: 0.0863 - acc: 0.9734 - val_loss: 0.0783 - val_acc: 0.9772

Epoch 7/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0803 - acc: 0.9758 - val_loss: 0.0704 - val_acc: 0.9790

Epoch 8/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0723 - acc: 0.9777 - val_loss: 0.0646 - val_acc: 0.9802

Epoch 9/20

60000/60000 [=====] - 8s 142us/step - loss: 0.0678 - acc: 0.9787 - val_loss: 0.0602 - val_acc: 0.9831

Epoch 10/20

60000/60000 [=====] - 8s 142us/step - loss: 0.0644 - acc: 0.9802 - val_loss: 0.0638 - val_acc: 0.9814

Epoch 11/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0583 - acc: 0.9819 - val_loss: 0.0568 - val_acc: 0.9834

Epoch 12/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0559 - acc: 0.9826 - val_loss: 0.0593 - val_acc: 0.9831

Epoch 13/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0529 - acc: 0.9838 - val_loss: 0.0670 - val_acc: 0.9802

Epoch 14/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0502 - acc: 0.9838 - val_loss: 0.0541 - val_acc: 0.9845

Epoch 15/20

60000/60000 [=====] - 8s 141us/step - loss: 0.0468 - acc: 0.9851 - val_loss: 0.0590 - val_acc: 0.9848

Epoch 16/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0438 - acc: 0.9863 - val_loss: 0.0633 - val_acc: 0.9835

Epoch 17/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0444 - acc: 0.9860 - val_loss: 0.0622 - val_acc: 0.9836

Epoch 18/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0423 - acc: 0.9872 - val_loss: 0.0568 - val_acc: 0.9834

Epoch 19/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0396 - acc: 0.9871 - val_loss: 0.0584 - val_acc: 0.9849

Epoch 20/20

60000/60000 [=====] - 9s 142us/step - loss: 0.0387 - acc: 0.9882 -

```
00000/00000 |> 100% 11240/step 10000.0/loss acc: 0.9832  
val_loss: 0.0600 - val_acc: 0.9832
```

In [132]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06003061625857372

Test accuracy: 0.9832



In [133]:

```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).



Model 2: MLP + Dropout (0.30)

- #layers: 5
- activation: ReLU
- Weight Initializer: He Normal
- Optimizer: ADAM

In [134]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal()))
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='relu', kernel_initializer=he_normal()))
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='relu', kernel_initializer=he_normal()))
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='relu', kernel_initializer=he_normal()))
model_relu.add(Dense(32, activation='relu', kernel_initializer=he_normal()))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_161 (Dense)	(None, 512)	401920
dropout_48 (Dropout)	(None, 512)	0
dense_162 (Dense)	(None, 256)	131328
dropout_49 (Dropout)	(None, 256)	0
dense_163 (Dense)	(None, 128)	32896

dropout_50 (Dropout)	(None, 128)	0
dense_164 (Dense)	(None, 64)	8256
dense_165 (Dense)	(None, 32)	2080
dense_166 (Dense)	(None, 10)	330
=====		
Total params: 576,810		
Trainable params: 576,810		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 194us/step - loss: 0.4457 - acc: 0.8628 - val_loss: 0.1451 - val_acc: 0.9572

Epoch 2/20

60000/60000 [=====] - 4s 74us/step - loss: 0.1724 - acc: 0.9502 - val_loss: 0.1117 - val_acc: 0.9678

Epoch 3/20

60000/60000 [=====] - 4s 74us/step - loss: 0.1319 - acc: 0.9628 - val_loss: 0.0957 - val_acc: 0.9693

Epoch 4/20

60000/60000 [=====] - 4s 75us/step - loss: 0.1059 - acc: 0.9694 - val_loss: 0.0755 - val_acc: 0.9762

Epoch 5/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0885 - acc: 0.9738 - val_loss: 0.0729 - val_acc: 0.9785

Epoch 6/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0804 - acc: 0.9756 - val_loss: 0.0725 - val_acc: 0.9783

Epoch 7/20

60000/60000 [=====] - 4s 75us/step - loss: 0.0689 - acc: 0.9796 - val_loss: 0.0730 - val_acc: 0.9782

Epoch 8/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0650 - acc: 0.9807 - val_loss: 0.0749 - val_acc: 0.9788

Epoch 9/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0606 - acc: 0.9816 - val_loss: 0.0718 - val_acc: 0.9794

Epoch 10/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0548 - acc: 0.9834 - val_loss: 0.0708 - val_acc: 0.9813

Epoch 11/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0544 - acc: 0.9832 - val_loss: 0.0816 - val_acc: 0.9791

Epoch 12/20

60000/60000 [=====] - 4s 75us/step - loss: 0.0469 - acc: 0.9856 - val_loss: 0.0659 - val_acc: 0.9825

Epoch 13/20

60000/60000 [=====] - 5s 76us/step - loss: 0.0457 - acc: 0.9866 - val_loss: 0.0730 - val_acc: 0.9813

Epoch 14/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0411 - acc: 0.9875 - val_loss: 0.0662 - val_acc: 0.9829

Epoch 15/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0410 - acc: 0.9877 - val_loss: 0.0734 - val_acc: 0.9827

Epoch 16/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0397 - acc: 0.9886 - val_loss: 0.0705 - val_acc: 0.9828

Epoch 17/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0360 - acc: 0.9892 - val_loss: 0.0742 - val_acc: 0.9814

Epoch 18/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0390 - acc: 0.9885 - val_loss: 0.0655 - val_acc: 0.9829

Epoch 19/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0309 - acc: 0.9904 - val_loss: 0.0689 - val_acc: 0.9833

Epoch 20/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0335 - acc: 0.9901 - val_loss: 0.0675 - val_acc: 0.9833

In [135]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06747149933629462

Test accuracy: 0.9833



In [136]:

```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Model 3: MLP + BatchNormalization + Dropout (0.40)

- #layers: 5
- activation: ReLU
- Weight Initializer: RandomNormal
- Optimizer: ADAM

In [137]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.050, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))
model_relu.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))
model_relu.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.088, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.4))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.176, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_167 (Dense)	(None, 512)	401920
batch_normalization_64 (Batch Normalization)	(None, 512)	2048
dropout_51 (Dropout)	(None, 512)	0
dense_168 (Dense)	(None, 256)	131328
batch_normalization_65 (Batch Normalization)	(None, 256)	1024

dropout_52 (Dropout)	(None, 256)	0
dense_169 (Dense)	(None, 128)	32896
batch_normalization_66 (Batch Normalization)	(None, 128)	512
dropout_53 (Dropout)	(None, 128)	0
dense_170 (Dense)	(None, 64)	8256
batch_normalization_67 (Batch Normalization)	(None, 64)	256
dense_171 (Dense)	(None, 32)	2080
batch_normalization_68 (Batch Normalization)	(None, 32)	128
dense_172 (Dense)	(None, 10)	330
=====		
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 291us/step - loss: 0.5734 - acc: 0.8261 - val_loss: 0.1564 - val_acc: 0.9542

Epoch 2/20

60000/60000 [=====] - 9s 145us/step - loss: 0.2228 - acc: 0.9344 - val_loss: 0.1220 - val_acc: 0.9634

Epoch 3/20

60000/60000 [=====] - 9s 153us/step - loss: 0.1739 - acc: 0.9496 - val_loss: 0.1009 - val_acc: 0.9690

Epoch 4/20

60000/60000 [=====] - 9s 148us/step - loss: 0.1416 - acc: 0.9580 - val_loss: 0.0903 - val_acc: 0.9737

Epoch 5/20

60000/60000 [=====] - 9s 145us/step - loss: 0.1279 - acc: 0.9622 - val_loss: 0.0767 - val_acc: 0.9766

Epoch 6/20

60000/60000 [=====] - 9s 145us/step - loss: 0.1174 - acc: 0.9654 - val_loss: 0.0806 - val_acc: 0.9762

Epoch 7/20

60000/60000 [=====] - 9s 145us/step - loss: 0.1058 - acc: 0.9683 - val_loss: 0.0778 - val_acc: 0.9760

Epoch 8/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0971 - acc: 0.9710 - val_loss: 0.0673 - val_acc: 0.9793

Epoch 9/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0893 - acc: 0.9730 - val_loss: 0.0737 - val_acc: 0.9785

Epoch 10/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0840 - acc: 0.9748 - val_loss: 0.0634 - val_acc: 0.9824

Epoch 11/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0816 - acc: 0.9757 - val_loss: 0.0611 - val_acc: 0.9824

Epoch 12/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0751 - acc: 0.9776 - val_loss: 0.0667 - val_acc: 0.9811

Epoch 13/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0715 - acc: 0.9783 - val_loss: 0.0640 - val_acc: 0.9810

Epoch 14/20

60000/60000 [=====] - 9s 143us/step - loss: 0.0658 - acc: 0.9798 - val_loss: 0.0632 - val_acc: 0.9824

Epoch 15/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0665 - acc: 0.9799 - val_loss: 0.0613 - val_acc: 0.9835

Epoch 16/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0623 - acc: 0.9809 - val_loss: 0.0619 - val_acc: 0.9819

Epoch 17/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0597 - acc: 0.9822 - val_loss: 0.0571 - val_acc: 0.9836

Epoch 18/20

60000/60000 [=====] - 9s 144us/step - loss: 0.0577 - acc: 0.9822 -

```

val_loss: 0.0557 - val_acc: 0.9836
Epoch 19/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0558 - acc: 0.9826 -
val_loss: 0.0590 - val_acc: 0.9833
Epoch 20/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0522 - acc: 0.9837 -
val_loss: 0.0556 - val_acc: 0.9836

```

In [138]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```


```

Test score: 0.055616749329864976
Test accuracy: 0.9836

```

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).

max_open_warning, RuntimeWarning)



In [139]:

```

w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```


```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



Model 4: MLP + BatchNormalization + Dropout (0.30)

- #layers: 5
- activation: sigmoid
- Weight Initializer: RandomNormal
- Optimizer: ADAM

In [140]:

```

model_relu = Sequential()
model_relu.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=Random
Normal(mean=0.0, stddev=0.039, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.
051, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.
072, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1
02, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1
44, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, valid

```

```
ation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_173 (Dense)	(None, 512)	401920
batch_normalization_69 (Batch Normalization)	(None, 512)	2048
dropout_54 (Dropout)	(None, 512)	0
dense_174 (Dense)	(None, 256)	131328
batch_normalization_70 (Batch Normalization)	(None, 256)	1024
dropout_55 (Dropout)	(None, 256)	0
dense_175 (Dense)	(None, 128)	32896
batch_normalization_71 (Batch Normalization)	(None, 128)	512
dropout_56 (Dropout)	(None, 128)	0
dense_176 (Dense)	(None, 64)	8256
batch_normalization_72 (Batch Normalization)	(None, 64)	256
dense_177 (Dense)	(None, 32)	2080
batch_normalization_73 (Batch Normalization)	(None, 32)	128
dense_178 (Dense)	(None, 10)	330
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 18s 303us/step - loss: 0.4027 - acc: 0.8826 - val_loss: 0.2217 - val_acc: 0.9337

Epoch 2/20

60000/60000 [=====] - 9s 147us/step - loss: 0.2438 - acc: 0.9262 - val_loss: 0.1567 - val_acc: 0.9532

Epoch 3/20

60000/60000 [=====] - 9s 147us/step - loss: 0.1918 - acc: 0.9430 - val_loss: 0.1241 - val_acc: 0.9624

Epoch 4/20

60000/60000 [=====] - 9s 146us/step - loss: 0.1679 - acc: 0.9493 - val_loss: 0.1181 - val_acc: 0.9643

Epoch 5/20

60000/60000 [=====] - 9s 146us/step - loss: 0.1455 - acc: 0.9564 - val_loss: 0.1064 - val_acc: 0.9688

Epoch 6/20

60000/60000 [=====] - 9s 147us/step - loss: 0.1307 - acc: 0.9600 - val_loss: 0.0982 - val_acc: 0.9722

Epoch 7/20

60000/60000 [=====] - 9s 147us/step - loss: 0.1157 - acc: 0.9651 - val_loss: 0.0866 - val_acc: 0.9719

Epoch 8/20

60000/60000 [=====] - 9s 147us/step - loss: 0.1076 - acc: 0.9669 - val_loss: 0.0824 - val_acc: 0.9745

Epoch 9/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0948 - acc: 0.9706 - val_loss: 0.0808 - val_acc: 0.9747

Epoch 10/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0923 - acc: 0.9713 - val_loss: 0.0797 - val_acc: 0.9771

Epoch 11/20

60000/60000 [=====] - 9s 146us/step - loss: 0.0816 - acc: 0.9751 - val_loss: 0.0727 - val_acc: 0.9778

Epoch 12/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0795 - acc: 0.9761 - val_loss: 0.0691 - val_acc: 0.9799

Epoch 13/20

60000/60000 [=====] - 9s 147us/step - loss: 0.0761 - acc: 0.9766 - val_loss: 0.0676 - val_acc: 0.9799

```

60000/60000 [=====] - 9s 14us/step - loss: 0.0761 - acc: 0.9766 -
val_loss: 0.0670 - val_acc: 0.9799
Epoch 14/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0684 - acc: 0.9792 -
val_loss: 0.0659 - val_acc: 0.9800
Epoch 15/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0652 - acc: 0.9793 -
val_loss: 0.0674 - val_acc: 0.9807
Epoch 16/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0631 - acc: 0.9801 - val_l
oss: 0.0662 - val_acc: 0.9812
Epoch 17/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0575 - acc: 0.9820 -
val_loss: 0.0635 - val_acc: 0.9819
Epoch 18/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0574 - acc: 0.9819 -
val_loss: 0.0643 - val_acc: 0.9805
Epoch 19/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0553 - acc: 0.9826 -
val_loss: 0.0682 - val_acc: 0.9813
Epoch 20/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0513 - acc: 0.9840 -
val_loss: 0.0617 - val_acc: 0.9820

```

In [141]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)


```


Test score: 0.06166264152140356

Test accuracy: 0.982

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).

max_open_warning, RuntimeWarning)





In [142]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Model 5: MLP + BatchNormalization + Dropout (0.30)

- #layers: 5
- activation: sigmoid
- Weight Initializer: RandomNormal
- Optimizer: adadelta

In [143]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.051, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.073, seed=None)))
```

```

0.12, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.102, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.144, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adadelta', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Layer (type)	Output Shape	Param #
dense_179 (Dense)	(None, 512)	401920
batch_normalization_74 (Batch Normalization)	(None, 512)	2048
dropout_57 (Dropout)	(None, 512)	0
dense_180 (Dense)	(None, 256)	131328
batch_normalization_75 (Batch Normalization)	(None, 256)	1024
dropout_58 (Dropout)	(None, 256)	0
dense_181 (Dense)	(None, 128)	32896
batch_normalization_76 (Batch Normalization)	(None, 128)	512
dropout_59 (Dropout)	(None, 128)	0
dense_182 (Dense)	(None, 64)	8256
batch_normalization_77 (Batch Normalization)	(None, 64)	256
dense_183 (Dense)	(None, 32)	2080
batch_normalization_78 (Batch Normalization)	(None, 32)	128
dense_184 (Dense)	(None, 10)	330
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

```

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 19s 316us/step - loss: 0.4114 - acc: 0.8797 - val_loss: 0.2380 - val_acc: 0.9287
Epoch 2/20
60000/60000 [=====] - 9s 156us/step - loss: 0.2509 - acc: 0.9246 - val_loss: 0.1653 - val_acc: 0.9520
Epoch 3/20
60000/60000 [=====] - 9s 157us/step - loss: 0.2036 - acc: 0.9381 - val_loss: 0.1451 - val_acc: 0.9578
Epoch 4/20
60000/60000 [=====] - 9s 154us/step - loss: 0.1751 - acc: 0.9468 - val_loss: 0.1324 - val_acc: 0.9632
Epoch 5/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1528 - acc: 0.9545 - val_loss: 0.1179 - val_acc: 0.9650
Epoch 6/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1415 - acc: 0.9574 - val_loss: 0.1063 - val_acc: 0.9698
Epoch 7/20
60000/60000 [=====] - 9s 154us/step - loss: 0.1299 - acc: 0.9610 - val_loss: 0.0977 - val_acc: 0.9719
Epoch 8/20

```

```

60000/60000 [=====] - 9s 155us/step - loss: 0.1196 - acc: 0.9635 -
val_loss: 0.0920 - val_acc: 0.9732
Epoch 9/20
60000/60000 [=====] - 9s 155us/step - loss: 0.1147 - acc: 0.9656 -
val_loss: 0.0883 - val_acc: 0.9744
Epoch 10/20
60000/60000 [=====] - 9s 155us/step - loss: 0.1083 - acc: 0.9668 -
val_loss: 0.0860 - val_acc: 0.9748
Epoch 11/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0996 - acc: 0.9697 -
val_loss: 0.0860 - val_acc: 0.9753
Epoch 12/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0948 - acc: 0.9703 -
val_loss: 0.0866 - val_acc: 0.9762
Epoch 13/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0896 - acc: 0.9730 -
val_loss: 0.0821 - val_acc: 0.9782
Epoch 14/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0857 - acc: 0.9733 -
val_loss: 0.0798 - val_acc: 0.9774
Epoch 15/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0828 - acc: 0.9744 -
val_loss: 0.0741 - val_acc: 0.9786
Epoch 16/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0790 - acc: 0.9758 -
val_loss: 0.0696 - val_acc: 0.9793
Epoch 17/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0752 - acc: 0.9759 -
val_loss: 0.0695 - val_acc: 0.9803
Epoch 18/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0721 - acc: 0.9767 -
val_loss: 0.0684 - val_acc: 0.9796
Epoch 19/20
60000/60000 [=====] - 9s 157us/step - loss: 0.0682 - acc: 0.9791 -
val_loss: 0.0729 - val_acc: 0.9788
Epoch 20/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0675 - acc: 0.9789 -
val_loss: 0.0698 - val_acc: 0.9804

```

In [144]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```


```

Test score: 0.06982594733461737
Test accuracy: 0.9804

```

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface

(matplotlib.pyplot.figure) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).
 max_open_warning, RuntimeWarning)



In [145]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Model 6: MLP + BatchNormalization + Dropout (0.30)

- #layers: 5
- activation: tanh
- Weight Initializer: glorot_normal

- weight_initializer: glorot_normal
- Optimizer: ADAM

In [146]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='tanh', input_shape=(input_dim,), kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(256, activation='tanh', kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(128, activation='tanh', kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.3))
model_relu.add(Dense(64, activation='tanh', kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dense(32, activation='tanh', kernel_initializer=glorot_normal()))
model_relu.add(BatchNormalization())
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_185 (Dense)	(None, 512)	401920
batch_normalization_79 (Batch Normalization)	(None, 512)	2048
dropout_60 (Dropout)	(None, 512)	0
dense_186 (Dense)	(None, 256)	131328
batch_normalization_80 (Batch Normalization)	(None, 256)	1024
dropout_61 (Dropout)	(None, 256)	0
dense_187 (Dense)	(None, 128)	32896
batch_normalization_81 (Batch Normalization)	(None, 128)	512
dropout_62 (Dropout)	(None, 128)	0
dense_188 (Dense)	(None, 64)	8256
batch_normalization_82 (Batch Normalization)	(None, 64)	256
dense_189 (Dense)	(None, 32)	2080
batch_normalization_83 (Batch Normalization)	(None, 32)	128
dense_190 (Dense)	(None, 10)	330
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 19s 312us/step - loss: 0.4344 - acc: 0.8725 - val_loss: 0.2088 - val_acc: 0.9388

Epoch 2/20

60000/60000 [=====] - 9s 145us/step - loss: 0.2380 - acc: 0.9296 - val_loss: 0.1496 - val_acc: 0.9580

Epoch 3/20

60000/60000 [=====] - 9s 146us/step - loss: 0.1820 - acc: 0.9455 - val_loss: 0.1261 - val_acc: 0.9636

Epoch 4/20

```

60000/60000 [=====] - 9s 158us/step - loss: 0.1496 - acc: 0.9563 -
val_loss: 0.1042 - val_acc: 0.9712
Epoch 5/20
60000/60000 [=====] - 9s 149us/step - loss: 0.1287 - acc: 0.9619 -
val_loss: 0.1016 - val_acc: 0.9722
Epoch 6/20
60000/60000 [=====] - 9s 146us/step - loss: 0.1195 - acc: 0.9648 -
val_loss: 0.0928 - val_acc: 0.9740
Epoch 7/20
60000/60000 [=====] - 9s 146us/step - loss: 0.1040 - acc: 0.9697 -
val_loss: 0.0841 - val_acc: 0.9737
Epoch 8/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0968 - acc: 0.9712 -
val_loss: 0.0780 - val_acc: 0.9763
Epoch 9/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0900 - acc: 0.9734 -
val_loss: 0.0769 - val_acc: 0.9776
Epoch 10/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0818 - acc: 0.9751 -
val_loss: 0.0716 - val_acc: 0.9796
Epoch 11/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0744 - acc: 0.9770 -
val_loss: 0.0713 - val_acc: 0.9801
Epoch 12/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0755 - acc: 0.9772 -
val_loss: 0.0691 - val_acc: 0.9818
Epoch 13/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0677 - acc: 0.9797 -
val_loss: 0.0741 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0613 - acc: 0.9815 -
val_loss: 0.0655 - val_acc: 0.9832
Epoch 15/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0602 - acc: 0.9821 -
val_loss: 0.0731 - val_acc: 0.9819
Epoch 16/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0596 - acc: 0.9817 -
val_loss: 0.0713 - val_acc: 0.9821
Epoch 17/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0548 - acc: 0.9834 -
val_loss: 0.0730 - val_acc: 0.9814
Epoch 18/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0562 - acc: 0.9825 -
val_loss: 0.0634 - val_acc: 0.9829
Epoch 19/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0508 - acc: 0.9847 -
val_loss: 0.0717 - val_acc: 0.9822
Epoch 20/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0498 - acc: 0.9849 -
val_loss: 0.0628 - val_acc: 0.9843

```

In [147]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06275533262640237
Test accuracy: 0.9843

/opt/conda/lib/python3.6/site-packages/matplotlib/pyplot.py:514: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).

max_open_warning, RuntimeWarning)



In [148]:


```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





Feedback Conclusion:

- I have used Kaggle platform to do this assignment as I found that kaggle is much much faster than Google colab.
- I have trained MLP Models with 2, 3 and 5 layers.
- I have used RandomNormal, He Normal and Glorot Normal weight initialization.
- I have used ReLU, sigmoid and tanh activation function.
- I have used AdaDelta and ADAM as optimizer.
- ADAM is faster than AdaDelta.
- I have also used BatchNormalization and Dropout.

In [149]:

```
table = PrettyTable()
table.field_names = ['Model #', 'Batch Normalization', 'Dropout + Value', 'Activation', 'Initialize
r', 'Optimizer', 'Accuracy']
table.add_row([1, "Yes", "Yes, 0.3", "ReLU", "RandomNormal", "ADAM", 0.9832])
table.add_row([2, "No", "Yes, 0.3", "ReLU", "He Normal", "ADAM", 0.9833])
table.add_row([3, "Yes", "Yes, 0.4", "ReLU", "RandomNormal", "ADAM", 0.9836])
table.add_row([4, "Yes", "Yes, 0.3", "sigmoid", "RandomNormal", "ADAM", 0.9820])
table.add_row([5, "Yes", "Yes, 0.3", "sigmoid", "RandomNormal", "AdaDelta", 0.9804])
table.add_row([6, "Yes", "Yes, 0.3", "tanh", "Glorot Normal", "ADAM", 0.9848])
print(table)
```

Model #	Batch Normalization	Dropout + Value	Activation	Initializer	Optimizer	Accuracy
1	Yes	Yes, 0.3	ReLU	RandomNormal	ADAM	0.983
2	No	Yes, 0.3	ReLU	He Normal	ADAM	0.983
3	Yes	Yes, 0.4	ReLU	RandomNormal	ADAM	0.983
4	Yes	Yes, 0.3	sigmoid	RandomNormal	ADAM	0.982
5	Yes	Yes, 0.3	sigmoid	RandomNormal	AdaDelta	0.980
6	Yes	Yes, 0.3	tanh	Glorot Normal	ADAM	0.984